

HR & Talent Transformation

Generative AI for Business Leaders & C-Suite | Cost Reduction & Efficiency

1. Why HR Is Ripe for AI Disruption

Human Resources sits at an unusual intersection: it handles some of the most strategic decisions an organisation makes (who to hire, how to develop talent, when to intervene on retention) yet spends most of its time on administrative processes that have barely changed in decades. Resume screening, interview scheduling, benefits enrollment, policy compliance, exit paperwork — these tasks consume an estimated 60–70 per cent of HR professionals' time, according to a 2024 Gartner survey, leaving precious little capacity for the strategic advisory work that CHROs aspire to deliver.

The mismatch is expensive. The Society for Human Resource Management (SHRM) estimates the average cost-per-hire at \$4,700, but for professional and technical roles, the all-in cost — including recruiter time, advertising, interviewing, onboarding, and first-year productivity ramp — ranges from \$15,000 to \$25,000. When your hiring process is slow, biased, or inconsistent, each bad hire or lost candidate compounds that cost. AI does not eliminate the need for human judgment in HR. It eliminates the administrative bottlenecks that prevent HR professionals from exercising that judgment where it matters most.

2. The Hiring Bottleneck — Anatomy of a Broken Process

The average corporate job posting attracts 250 applications. For popular roles at well-known employers, the number can reach 500–1,000. A recruiter screening those manually faces an impossible task: reading each resume takes 6–8 minutes for a thorough review, meaning 250 applications consume 25–33 hours of focused reading. In practice, recruiters spend 6–7 seconds per resume in an initial scan — just enough to form a first impression based on formatting, school name, company name, and job title. Decades of research confirm what this implies.

Real-World Incident — National Bureau of Economic Research (NBER) Audit Study (2021): Researchers sent 83,000 fictitious job applications to Fortune 500 companies, varying only the perceived race and ethnicity of applicant names. They found that applicants with distinctively white-sounding names received 9.5 per cent more callbacks than equally qualified applicants with distinctively Black-sounding names. Some companies showed no bias; others showed callback gaps as large as 24 per cent. The study demonstrated that unconscious bias in resume screening is measurable, persistent, and affects even companies with explicit diversity commitments.

The problem extends beyond bias. Fatigue effects mean that a resume reviewed at 9 AM receives different treatment than the same resume reviewed at 4 PM after 200 others. Format bias causes recruiters to favour polished layouts over content. Name-brand bias leads to over-indexing on prestigious university or employer names. And speed bias causes most applications to be rejected in seconds, meaning genuinely strong candidates with non-traditional backgrounds are routinely overlooked. AI screening does not have bad afternoons, does not recognise school prestige, and does not fatigue after the 200th resume.

3. AI Across the Talent Lifecycle — A Detailed Walkthrough

Intelligent Resume Screening and Candidate Matching

Modern AI screening tools (HireVue, Pymetrics, Eightfold.ai, Phenom) go beyond keyword matching. They build multi-dimensional candidate profiles based on skills, experience patterns, career trajectory, and inferred capabilities. A candidate who spent five years in supply chain logistics at a mid-size manufacturer may not match the exact keyword list for a 'supply chain analyst' posting, but the AI recognises the functional overlap and scores them appropriately. Companies using AI screening report 75 per cent time reduction in the initial screening phase and 50 per cent improvement in candidate quality reaching the interview stage — not because AI is smarter than recruiters, but because it is more consistent and does not tire.

Interview Scheduling and Candidate Communication

Interview scheduling is one of the most underestimated time sinks in recruitment. Coordinating availability across 3–5 interviewers, managing time zones, handling reschedules, sending reminders, and fielding candidate questions about logistics can consume 3–5 hours per hire. AI scheduling assistants (Paradox's Olivia, GoodTime, Calendly's AI features) handle the entire coordination cycle: they propose times based on everyone's calendar availability, send invitations, confirm attendance, and automatically reschedule when conflicts arise. For high-volume hiring — retail, hospitality, contact centres — this automation is transformative, converting days of scheduling work into minutes.

Personalised Learning and Development

Traditional corporate training follows a one-size-fits-all model: everyone in a role takes the same courses, regardless of their existing skills or preferred learning style. The result is predictable — experts are bored, novices are overwhelmed, and completion rates rarely exceed 30–40 per cent. AI-powered L&D platforms (Cornerstone, Degreed, Docebo, LinkedIn Learning with AI recommendations) assess each employee's current skills through assessments and work output analysis, then create personalised learning pathways. The AI adapts based on engagement: if an employee struggles with a module, it offers simpler prerequisites; if they breeze through, it accelerates to advanced content. Organisations using AI-personalised learning report 40–60 per cent higher course completion rates and measurably faster skill acquisition.

Performance Management Assistance

Performance reviews are universally dreaded by managers and employees alike, partly because they are time-consuming and partly because they are inconsistent. Manager A writes detailed, balanced reviews; Manager B writes two sentences. AI can analyse 360-degree feedback data, goal-achievement records, project outcomes, and peer assessments to generate draft performance reviews that managers then edit and personalise. This ensures a consistent baseline across the organisation while still preserving the manager's voice and judgment. The time saving is significant — a 2024 Lattice survey found that managers spend an average of 17 hours per review cycle on performance documentation. AI can reduce that to 4–6 hours.

Engagement Analysis and Retention Prediction

Losing an employee is expensive. SHRM estimates the replacement cost at 50–200 per cent of annual salary, depending on the role. For a senior engineer earning \$150,000, replacement cost can exceed \$300,000 when you account for recruiting, onboarding, lost productivity during the vacancy, and the ramp-up time for the new hire. AI retention tools (Culture Amp, Glint, Workday People Analytics, Visier) analyse patterns in engagement survey responses, communication frequency, meeting attendance, project participation, and other behavioural signals to predict flight risk. The goal is not surveillance — it is early intervention. If an employee's engagement pattern shifts in ways that historically correlate with departure, their manager or HR partner is alerted and can have a proactive conversation about career growth, workload, or compensation.

4. The Bias Question — AI as Both Problem and Solution

The claim that 'AI removes bias from hiring' requires nuance. AI trained on biased historical data will reproduce those biases — it learns patterns from whatever data it is given. If your historical hiring data shows that men were promoted faster than women, an AI trained on that data may learn to favour male candidates.

Real-World Incident — Amazon's AI Recruiting Tool (2018): Amazon built an AI recruiting system trained on 10 years of hiring data. Because Amazon's tech workforce was predominantly male during that period, the AI learned to penalise resumes containing the word 'women's' (as in 'women's chess club captain') and downgrade graduates of all-women's colleges. Amazon detected the bias during internal testing and scrapped the tool before external deployment. The incident became a cautionary tale about training data quality and a reminder that AI reflects the patterns in its data, for better or worse.

However, when designed correctly, AI can be significantly less biased than human screening. The key requirements are: training on debiased data that represents the diversity you want to achieve, regular bias audits that test AI outputs across demographic groups, skills-based assessment rather than proxy-based screening (evaluating what a candidate can do, not where they went to school), and transparent scoring criteria that can be reviewed and challenged. Unilever's AI hiring system — covered in the next lecture — demonstrates how AI, properly implemented, can increase diversity while improving hire quality.

5. Cost-Benefit Analysis — HR AI by the Numbers

AI Application	Typical Investment	Typical Annual Savings/Value	Payback Period
Resume screening & matching	\$30–\$80/employee/year	40–60% reduction in time-to-hire; \$5,000–\$10,000 saved per hire	1–3 months
Interview scheduling automation	\$5–\$15/employee/year	3–5 hours saved per hire in coordination time	Immediate
Personalised L&D	\$20–\$60/employee/year	40–60% higher course completion; faster skill acquisition	3–6 months
Performance review	\$10–	10–13 hours saved per	1–2 cycles

assistance	\$25/employee/year	manager per review cycle	
Retention prediction	\$15– \$40/employee/year	Preventing one departure saves 50–200% of that salary	1 prevented departure

Key Point: The total investment for a comprehensive HR AI stack is typically \$80–\$220 per employee per year. For a 500-person company, that is \$40,000–\$110,000 annually. If AI prevents just 5 voluntary departures (at a conservative \$75,000 replacement cost each), the retention value alone — \$375,000 — covers the entire investment 3–9 times over. All other savings (faster hiring, better training, reduced admin) are pure additional return.

6. The Three-Phase Implementation Roadmap

HR AI implementation follows a natural progression from operational automation (immediate savings) to strategic capability (long-term competitive advantage).

Phase	Focus	Timeline	Key Actions
Phase 1	Recruitment automation	Months 1–3	Deploy AI screening; automate scheduling; implement candidate communication bot
Phase 2	Learning & development	Months 3–6	Launch personalised learning paths; integrate with performance data; build skills taxonomy
Phase 3	Strategic workforce intelligence	Months 6–12	Add retention prediction; enable AI-assisted performance reviews; build workforce planning models

Phase 1 delivers the fastest ROI and builds organisational confidence. Start here regardless of your overall ambition. Phase 2 creates lasting capability by ensuring employees develop the skills the organisation needs. Phase 3 turns HR into a predictive, strategic function rather than a reactive, administrative one.

7. The Leaky Bucket Analogy — Why Retention Outweighs Recruitment

Imagine your organisation as a bucket of water. Recruitment is the tap filling the bucket. Retention is the integrity of the bucket itself. Most organisations invest heavily in the tap — employer branding, job boards, recruiters, signing bonuses — while ignoring the holes in the bucket. Every departed employee is water leaking out. AI recruitment tools make the tap flow faster. AI retention tools patch the holes. The maths always favours patching holes: retaining one senior employee saves more than hiring two new ones when you account for institutional knowledge loss, team disruption, and ramp-up time.

Gallup estimates that replacing a highly skilled employee costs 200 per cent of their annual salary when all factors are included. For a team of 50 knowledge workers with 15 per cent annual turnover, that is 7–8 departures per year, costing \$1–\$2 million in replacement and productivity losses. Reducing turnover by even 3–4 percentage points through AI-powered early intervention translates to 1.5–2 fewer departures per year, saving \$300,000–\$600,000 annually. The AI tool that enables this costs perhaps \$10,000–\$20,000 for the same 50-person team.

8. Ethical Guardrails for HR AI

HR AI operates on some of the most sensitive data an organisation handles — personal information, performance assessments, compensation data, health-related accommodations. The ethical stakes are proportionally high.

- **Transparency:** Candidates and employees should know when AI is used in decisions that affect them. Many jurisdictions now require this by law (New York City's Local Law 144, the EU AI Act's high-risk classification for employment AI).
- **Bias audits:** Run quarterly disparate-impact analyses across gender, race, age, and disability status. If AI screening produces materially different outcomes for protected groups, investigate and correct before continuing.
- **Human final decisions:** AI should inform, not determine, hiring and termination decisions. A human must review and approve any consequential personnel action.
- **Data minimisation:** Collect and process only the data necessary for the AI's purpose. Social media monitoring, health data analysis, and personal relationship mapping cross ethical boundaries even if technically possible.
- **Opt-out pathways:** Candidates should have the option to request human-only assessment if they object to AI evaluation. This is a legal requirement under some regulations and a trust-building practice everywhere.

9. Key Terms Glossary

Term	Definition
Cost-per-hire	Total recruitment cost (advertising, recruiter time, tools, interviewing, onboarding) divided by number of hires — SHRM average is ~\$4,700 but varies widely by role seniority
Unconscious bias	Systematic patterns of deviation from rationality in judgment, where individuals unknowingly favour or disfavour candidates based on irrelevant characteristics (name, school, gender, age)
Skills-based hiring	Evaluating candidates on demonstrated abilities and competencies rather than credentials, pedigree, or proxy indicators — increasingly enabled by AI assessment tools
Disparate impact	A legal standard where a facially neutral practice (like an AI screening algorithm) produces significantly different outcomes for a protected group — triggers regulatory scrutiny even without discriminatory intent
Flight risk	The probability that an employee will voluntarily leave the organisation within a defined period — AI retention tools estimate this based on engagement, behavioural, and market signals
Neuroscience-based assessment	Candidate evaluation using cognitive games that measure traits like risk tolerance, attention, and problem-solving speed — used by platforms like Pymetrics as an alternative to resume screening
360-degree feedback	Performance evaluation that incorporates input from an employee's manager, peers, direct reports, and sometimes customers — AI can synthesise these diverse inputs into coherent draft reviews
Local Law 144 (NYC)	New York City regulation effective July 2023 requiring employers to

	conduct annual bias audits of automated employment decision tools and notify candidates when AI is used in hiring
--	---

10. Quick Revision Checklist

- HR professionals spend 60–70% of their time on administrative tasks; AI targets this administrative burden, not HR's strategic role
- The average corporate posting attracts 250 applications; manual screening takes 25–33 hours and is plagued by fatigue, format, and name-brand bias
- AI screening delivers 75% time reduction in initial screening and 50% improvement in candidate quality reaching interviews
- The five key AI applications in HR: resume screening, scheduling automation, personalised L&D, performance review assistance, and retention prediction
- Amazon's 2018 AI recruiting failure demonstrated that AI trained on biased data reproduces that bias — training data quality is paramount
- The NBER 2021 study found 9.5% callback gap based solely on perceived race of applicant names — AI can reduce this if properly debiased
- Total HR AI investment typically runs \$80–\$220 per employee per year; preventing just a few voluntary departures covers the entire cost
- Three-phase implementation: recruitment automation (months 1–3) → learning & development (months 3–6) → strategic workforce intelligence (months 6–12)
- Ethical guardrails are mandatory: transparency, bias audits, human final decisions, data minimisation, and opt-out pathways
- Replacing a highly skilled employee costs 200% of annual salary (Gallup) — retention ROI almost always exceeds recruitment ROI

20 Exam-Based MCQs — HR & Talent Transformation

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. Scenario: A fast-growing tech company receives 800 applications for a senior software engineer position. The HR team of 3 recruiters needs to screen all applications and begin first-round interviews within 2 weeks. Each thorough resume review takes 7 minutes. What is the approximate total screening time if done manually, and what does this imply for the team?

- A) 56 hours total — manageable within 2 weeks alongside other duties
- B) 93 hours total — each recruiter would need to spend 31 hours on screening alone, consuming most of their 2-week capacity
- C) 40 hours total — one recruiter can handle it in a single work week
- D) 120 hours total — impossible without hiring temporary staff

Answer: B **Explanation:** $800 \times 7 \text{ minutes} = 5,600 \text{ minutes} = 93.3 \text{ hours}$. Divided among 3 recruiters = 31.1 hours each over 2 weeks (80 working hours), consuming 39% of their total capacity on a single requisition.

A and C understate the time. D overstates it.

Q2. The NBER 2021 audit study found that applicants with distinctively white-sounding names received what percentage more callbacks than equally qualified applicants with distinctively Black-sounding names?

- A) 2.5%
- B) 5.0%
- C) 9.5%
- D) 18%

Answer: C **Explanation:** C is correct — the study found a 9.5% callback gap based solely on name perception, with some companies showing gaps as large as 24%. A and B understate the finding. D overstates the average, though some individual companies reached that level.

Q3. *Scenario:* A financial services firm is implementing AI resume screening. The CHRO wants to ensure the tool does not replicate historical biases. The firm's workforce is currently 72% male in investment roles due to decades of industry-wide gender imbalance. What is the MOST effective approach to prevent gender bias in the AI screening tool?

- A) Train the AI exclusively on resumes of successful female employees to correct the imbalance
- B) Remove gender-identifying information from training data and use skills-based assessment criteria rather than historical hiring patterns
- C) Set a 50/50 gender quota in the AI's output regardless of qualifications
- D) Use the AI only for male candidates and screen female candidates manually to ensure fairness

Answer: B **Explanation:** B addresses the root cause by removing proxies for gender and evaluating on objective skills. A creates reverse bias and reduces the training data quality. C imposes quotas that may conflict with merit and legal requirements. D is discriminatory in its own right and defeats the purpose of AI efficiency.

Q4. Amazon's 2018 AI recruiting tool failure is BEST explained by:

- A) A software bug that randomly penalised certain candidates
- B) The AI was trained on 10 years of predominantly male hiring data and learned to penalise indicators of female gender
- C) External hackers manipulated the AI's scoring algorithm
- D) The AI was intentionally designed to favour male candidates

Answer: B **Explanation:** B is correct. The AI learned patterns from historical data where men were disproportionately hired and promoted in technical roles. It inferred that female indicators correlated with lower success probability. A, C, and D are factually wrong — the issue was data bias, not bugs, attacks, or intentional design.

Q5. Scenario: A 500-person manufacturing company spends \$80 per employee per year on an HR AI platform covering recruitment and retention. Average employee salary is \$65,000. Current annual voluntary turnover is 18% (90 departures). After 12 months with AI, turnover drops to 14% (70 departures). SHRM estimates replacement cost at 50% of salary. What is the approximate net annual value of the retention improvement?

- A) \$210,000
- B) \$610,000
- C) \$650,000
- D) \$1,300,000

Answer: B **Explanation:** Departures prevented: $90 - 70 = 20$. Replacement cost per departure: $\$65,000 \times 0.50 = \$32,500$. Savings: $20 \times \$32,500 = \$650,000$. AI cost: $500 \times \$80 = \$40,000$. Net value: $\$650,000 - \$40,000 = \$610,000$. A understates the savings. C is gross savings before AI cost. D uses 100% rather than 50% replacement cost.

Q6. Which phase of HR AI implementation should be deployed FIRST for fastest ROI?

- A) Strategic workforce planning
- B) AI-assisted performance reviews
- C) Recruitment automation (screening, scheduling, candidate communication)
- D) Employee engagement analysis

Answer: C **Explanation:** C delivers the fastest measurable ROI because recruitment has clear costs (time-to-hire, cost-per-hire) and immediate efficiency gains. A and D require data accumulation over months. B requires a full performance cycle to demonstrate value.

Q7. Scenario: A retail chain's HR director wants to deploy AI for interview scheduling across 200 store locations. Each location hires 5–10 employees per month. Currently, store managers spend an average of 45 minutes per candidate coordinating interview times via phone and email. What is the approximate monthly time saved across all locations if AI scheduling reduces coordination time by 85%?

- A) 500 hours
- B) 950 hours
- C) 1,900 hours
- D) 3,800 hours

Answer: B **Explanation:** Average hires: $200 \text{ locations} \times 7.5 \text{ avg candidates} = 1,500 \text{ candidates/month}$. Current time: $1,500 \times 45 \text{ min} = 67,500 \text{ min} = 1,125 \text{ hours}$. 85% reduction = 956 hours saved ≈ 950 . A understates by using fewer locations or lower candidate count. C and D overstate by using higher multipliers.

Q8. Personalised AI learning platforms typically improve course completion rates by:

- A) 10–15%
- B) 20–30%

- C) 40–60%
- D) 80–90%

Answer: C **Explanation:** C is correct. The improvement comes from matching content difficulty and style to individual learner needs, eliminating boredom for experts and overwhelm for novices. A and B understate the typical improvement. D overstates it — personalisation helps significantly but does not make all courses compelling.

Q9. Scenario: An HR team has deployed an AI retention prediction tool. The tool flags a high-performing engineering manager as a flight risk based on declining meeting attendance, reduced Slack activity, and a pattern that historically correlates with departures within 90 days. The engineering VP is sceptical and says 'She just got promoted — there's no way she's leaving.' What is the MOST appropriate action?

- A) Ignore the AI prediction — the VP knows their team better than an algorithm
- B) Immediately offer a retention bonus to prevent the departure
- C) Share the AI's finding with the VP and suggest a casual check-in conversation about career satisfaction, workload, and growth opportunities — the goal is early intelligence, not certainty
- D) Place the employee under closer monitoring to gather more data before acting

Answer: C **Explanation:** C uses the AI's prediction as an early warning signal that prompts a low-risk, high-value conversation. A dismisses useful data. B is premature and potentially insulting. D delays action and shifts into surveillance territory.

Q10. Under New York City's Local Law 144, employers using automated employment decision tools must:

- A) Obtain candidate consent before using any technology in the hiring process
- B) Conduct annual bias audits and notify candidates when AI is used in hiring or promotion decisions
- C) Use only AI tools approved by the NYC Department of Labor
- D) Ensure AI tools are developed by NYC-based vendors

Answer: B **Explanation:** B correctly states the two core requirements of Local Law 144: annual independent bias audits and candidate notification. A is broader than the law requires. C and D are fabricated requirements.

Q11. Gallup estimates that replacing a highly skilled employee costs approximately what percentage of their annual salary?

- A) 25%
- B) 50%
- C) 100%
- D) 200%

Answer: D **Explanation:** D is correct. Gallup's estimate includes recruiting costs, lost productivity during

vacancy, onboarding, training, and the ramp-up period before the new hire reaches full productivity. Lower estimates (50–100%) often exclude institutional knowledge loss and team disruption costs.

Q12. *Scenario:* A healthcare organisation is implementing AI for nurse recruitment. The compliance officer raises a concern: the AI assessment platform stores candidate data on cloud servers. Some candidate data includes health-related accommodation requests. The organisation is subject to HIPAA and state privacy laws. What is the MOST appropriate step before proceeding?

- A) Proceed — HIPAA only applies to patient data, not employee data
- B) Ensure the AI vendor has appropriate data security certifications, sign a data processing agreement, and exclude health-related data from AI processing
- C) Store all data on-premises to avoid cloud security concerns
- D) Use the AI only for candidates who consent to cloud data storage

Answer: B **Explanation:** B addresses the specific risks: vendor security, contractual data protection, and data minimisation (excluding sensitive health data from AI processing). A is a dangerous oversimplification — while HIPAA primarily covers patient data, state privacy laws and ADA requirements apply to employee health information. C is unnecessarily restrictive. D does not address the underlying data protection requirements.

Q13. The 'leaky bucket' analogy in HR AI refers to:

- A) Data leaks in AI systems processing employee information
- B) Organisations investing heavily in recruitment (filling the bucket) while ignoring retention (the holes), when fixing retention delivers higher ROI
- C) Candidates dropping out of lengthy hiring processes
- D) Budget overruns in HR technology implementations

Answer: B **Explanation:** B captures the analogy: recruitment is the tap, retention is the bucket integrity. Most organisations over-invest in the tap while ignoring the holes. AI retention tools address the holes — which is mathematically more valuable than faster filling.

Q14. *Scenario:* A company's AI screening tool has been in production for 6 months. A quarterly bias audit reveals that candidates over 50 are being screened out at a rate 22% higher than younger candidates with equivalent qualifications. What should the HR team do FIRST?

- A) Continue using the tool while investigating — stopping it would slow hiring
- B) Immediately pause AI screening for all candidates over 50, investigate the training data for age-correlated proxies, and correct before resuming
- C) Add age as a protected variable that the AI cannot see
- D) Report the finding to legal but continue operations unchanged

Answer: B **Explanation:** B follows the correct sequence: stop the harm, investigate the cause, fix before resuming. The 22% disparate impact is legally significant and ethically urgent. A and D allow ongoing discrimination. C is a partial fix — age proxies (graduation year, years of experience format) can recreate

the bias even if age itself is hidden.

Q15. Which of the following is the BEST example of 'skills-based hiring' enabled by AI?

- A) Requiring all candidates to hold a bachelor's degree from an accredited university
- B) Using neuroscience-based cognitive games to assess problem-solving, learning agility, and attention — independent of educational background
- C) Screening for candidates who have worked at Fortune 500 companies
- D) Prioritising candidates with the longest tenure at previous employers

Answer: B **Explanation:** B directly measures capabilities rather than proxy indicators. A, C, and D all use proxies (degree, employer brand, tenure) that correlate weakly with actual job performance and introduce bias against non-traditional candidates.

Q16. Scenario: A Lattice survey found that managers spend an average of 17 hours per review cycle on performance documentation. A company with 40 managers conducts reviews twice per year. If AI reduces documentation time by 70%, how many total manager-hours are saved annually?

- A) 476 hours
- B) 952 hours
- C) 1,360 hours
- D) 2,380 hours

Answer: B **Explanation:** Total current hours: 40 managers × 17 hours × 2 cycles = 1,360 hours/year. 70% savings = 1,360 × 0.70 = 952 hours. A calculates for one cycle only. C is the total current hours, not the savings. D uses a higher multiplier.

Q17. What is the PRIMARY ethical concern with AI-powered employee engagement monitoring?

- A) The technology is too expensive for most organisations
- B) The boundary between early-warning intelligence and employee surveillance is unclear, risking trust erosion if perceived as monitoring rather than supporting
- C) Engagement data is too subjective to be useful
- D) Only large enterprises have enough data for accurate predictions

Answer: B **Explanation:** B identifies the core tension: tools that analyse Slack activity, meeting attendance, and email patterns for retention prediction can feel like Big Brother if not communicated transparently. A, C, and D are secondary or inaccurate concerns.

Q18. According to SHRM, the average cost-per-hire across all industries is approximately:

- A) \$1,500
- B) \$4,700
- C) \$12,000

D) \$25,000

Answer: B **Explanation:** B is correct per SHRM's published data. A is too low even for entry-level roles. C and D represent all-in costs for professional/technical and executive roles respectively, not the cross-industry average.

Q19. Scenario: *A company's AI L&D platform recommends that a marketing manager take a data analytics course based on a skills gap analysis. The marketing manager objects, saying 'Data analytics isn't part of my job description.' The L&D director reviews the recommendation and finds it is based on a company-wide initiative to build data literacy across all departments. How should the L&D director handle this situation?*

- A) Override the employee's objection — the AI's recommendation aligns with company strategy
- B) Explain the strategic context (company-wide data literacy initiative), show how the skill connects to the marketing manager's career growth, and offer flexibility in timing or course format
- C) Remove the recommendation — employees should choose their own development path
- D) Escalate to the marketing manager's supervisor to mandate the training

Answer: B **Explanation:** B combines transparency (explaining the why), relevance (connecting to career growth), and autonomy (offering flexibility). A and D impose without context, which breeds resentment. C abandons strategic skill-building.

Q20. Which statement BEST summarises the relationship between AI and HR professionals?

- A) AI will replace most HR roles within 5 years
- B) AI handles administrative and data-processing tasks so HR professionals can focus on strategic advisory, relationship management, and complex judgment calls
- C) AI is only useful for large HR departments with 50+ employees
- D) AI eliminates the need for HR policies and governance frameworks

Answer: B **Explanation:** B captures the augmentation model: AI removes the administrative burden (60–70% of current HR time) so professionals can focus on the high-value, human-centric work that actually drives organisational outcomes. A overstates displacement. C ignores that even small HR teams benefit from automation. D is backwards — AI increases the need for governance frameworks.