

WORKSHOP ACTIVITY: Cost-Saving Calculator

Generative AI for Business Leaders & C-Suite | Cost Reduction & Efficiency

1. Why Quantification Changes Everything

AI conversations in boardrooms often stall at the same point: everyone agrees the technology is promising, but nobody can put a number on what it means for this specific organisation. The cost-saving calculator workshop bridges that gap. It forces leaders to move from abstract enthusiasm ('AI could save us money') to concrete projections ('AI could save us \$340,000 in the next twelve months on document processing alone').

This matters because investment decisions are made on numbers, not narratives. A CFO will not approve a \$50,000 annual AI tool subscription based on a McKinsey article about industry trends. They will approve it based on a credible analysis showing a 500% return on that \$50,000. This workshop gives you the framework to produce that analysis using your own organisation's real data.

The approach is deliberately conservative. The calculator uses a 40% efficiency gain assumption — well below the 60–85% improvements documented in case studies like Klarna, JPMorgan, and Walmart. Conservative projections build credibility with sceptical finance teams and leave room for upside surprise.

2. The Three Task Categories: A Complete Framework

The workshop identifies three categories of repetitive work that AI handles well. Understanding each category deeply — and honestly assessing your organisation's time allocation — is the foundation of an accurate projection.

Category 1: Document Processing

This encompasses any work that involves reading, reviewing, creating, extracting data from, or transforming documents. The breadth of this category surprises most leaders. It includes contract review, invoice processing, report generation, regulatory filing preparation, proposal writing, presentation creation, policy drafting, form processing, and email-to-action conversion. In a typical knowledge-work organisation, document processing consumes 20–35% of total professional time, according to a 2023 McKinsey Global Institute analysis of workplace productivity.

Category 2: Routine Communication

This includes customer service inquiries, internal status updates, meeting scheduling, standard email responses, FAQ handling, ticket triage, and any communication that follows predictable patterns. Customer-facing teams often spend 40–60% of their time on communications that fall into recognisable patterns — inquiries that have been answered before, updates that follow a template, or scheduling exchanges that could be automated.

Category 3: Data Work

Data entry, report compilation, spreadsheet analysis, system updates, data validation, reconciliation, dashboard creation, and any task that involves moving data between systems or transforming it into

actionable formats. Finance teams, operations teams, and analytics functions are particularly heavy in this category, often spending 30–50% of their time on data manipulation that follows rules a machine could apply.

Key Point: Most leaders underestimate the time their teams spend on repetitive tasks by 30–40%. The reason is normalisation — when everyone has always spent five hours a week on report generation, it stops being visible as waste. The honest audit this exercise requires often produces the workshop's first 'eye-opening' moment.

3. Calculating Fully Loaded Employee Cost

The accuracy of your savings projection depends on using realistic costs. Using base salary alone understates the true cost by 25–40%. The fully loaded cost includes every expense the organisation incurs for each employee hour.

Component	Typical Range	How to Calculate
Base salary	100% of salary	Annual salary figure from HR/finance
Benefits (health, retirement, insurance)	20–35% of salary	Total benefits cost / total employees, or ask HR for the per-employee figure
Payroll taxes	7–10% of salary	Employer-side FICA, state unemployment, workers' comp
Overhead (office space, equipment, IT)	10–20% of salary	Total overhead / total employees — varies significantly by industry
Fully loaded total	137–165% of salary	Sum of all components

The Quick Formula

If you cannot get precise figures, use the rule of thumb from the workshop: multiply annual salary by 1.4 to approximate fully loaded cost, then divide by 2,000 (standard work hours per year) for the hourly rate. For an employee earning \$70,000 per year: $\$70,000 \times 1.4 = \$98,000$ fully loaded / 2,000 hours = \$49 per hour. This approximation is close enough for initial projections. Refine with actual figures when presenting to finance for budget approval.

4. The Savings Formula: Step by Step

The complete calculation follows a four-step sequence. Each step builds on the previous one.

Step	Formula	Example
1. Annual hours on repetitive tasks	Weekly hours x 52 weeks	50 hrs/week x 52 = 2,600 hrs/year
2. Annual cost of repetitive tasks	Annual hours x fully loaded hourly rate	2,600 hrs x \$49/hr = \$127,400
3. Potential savings (40% efficiency gain)	Annual cost x 0.40	\$127,400 x 0.40 = \$50,960
4. Net savings after	Potential savings minus annual AI	\$50,960 - \$19,200 (20 users x \$80/month x

AI tool cost	tool cost	12) = \$31,760
--------------	-----------	----------------

Apply this calculation to each of your three task categories, then sum the results for your total organisational opportunity. Most leaders find that the combined savings across all three categories reaches six figures, even for modestly sized teams.

5. Understanding AI Tool Costs: The \$80/User/Month Benchmark

The workshop uses \$80 per user per month as a realistic benchmark for enterprise AI tools. This figure deserves context.

Tool Category	Examples	Typical Cost per User/Month	Best For
General-purpose enterprise AI	ChatGPT Enterprise, Microsoft 365 Copilot, Google Gemini for Workspace	\$25–\$60	Broad productivity gains across all three task categories
Specialised departmental AI	Legal AI (Kira, Luminance), Sales AI (Gong, Clari), HR AI (Eightfold, Phenom)	\$50–\$200	Deep capability in one specific function
Custom AI solutions	Fine-tuned models, internal platforms, API-based workflows	\$10–\$100+ (volume-based)	Organisations with unique data or workflows needing tailored solutions

The \$80 benchmark is conservative for general-purpose tools and may be low for specialised platforms. When building your business case, get actual vendor quotes for a more precise calculation. Many vendors offer volume discounts, so the per-user cost decreases as you scale across the organisation.

6. ROI and Payback Period: The Metrics That Move CFOs

Two metrics drive investment decisions: ROI (how much you get back for what you spend) and payback period (how quickly the investment pays for itself).

ROI Calculation

ROI = (Annual savings / Annual AI tool cost) x 100. Using the workshop example: \$100,800 net savings / \$19,200 annual tool cost x 100 = 525% ROI. For every dollar invested, you get \$5.25 back. By comparison, most corporate investments target 15–25% ROI. A 525% return puts AI tool investment in a category that makes traditional capital projects look anemic.

Payback Period

Payback period = Annual tool cost / Monthly savings. Using the workshop example: \$19,200 annual cost / (\$100,800 / 12 months) = \$19,200 / \$8,400 = 2.3 months. Your investment pays for itself in roughly ten weeks. Compare this to typical enterprise software deployments with 18–36 month payback periods, or physical infrastructure investments that may take 3–5 years to return.

Real-World Incident — Deloitte AI ROI Survey (2023): Deloitte surveyed 2,620 executives globally about their AI investments. Among organisations that had deployed AI at scale, the median reported ROI was 3.5x (350%), with the top quartile reporting over 5x (500%). Critically, 79% of respondents said their AI investments had already paid for themselves within the first year. The survey also found that organisations which started with a clear, quantified business case — like the one this workshop produces — were 2.4 times more likely to achieve their target ROI than those that deployed AI without specific financial targets.

7. From Calculator to Business Case: Making It Actionable

The workshop's final exercise asks you to commit to a priority and identify the decision-maker. This step transforms a spreadsheet exercise into an organisational action. A rigorous business case based on your calculator results includes five elements.

- **The problem quantified.** *'Our team spends 2,600 hours per year on document processing at a cost of \$127,400. This work is repetitive, error-prone, and prevents our professionals from focusing on strategic activities.'*
- **The solution proposed.** *'Deploy [specific AI tool] across [specific team/function] to automate [specific task category]. Expected efficiency gain: 40% (conservative, based on documented case studies).'*
- **The financial impact.** *'Annual savings: \$50,960. Tool cost: \$19,200/year. Net savings: \$31,760. ROI: 265%. Payback period: 4.5 months.'*
- **The risk mitigation.** *'Pilot deployment on [one specific process] for 60 days. Success criteria: measurable time savings validated by the team. Expansion only after pilot results confirmed.'*
- **The ask.** *'Approve \$19,200 annual subscription and 2 hours/week of project sponsor time for 8 weeks to manage the pilot.'*

This format works because it speaks the language of financial decision-makers: quantified problem, specific solution, measurable return, controlled risk, and a clear ask. The calculator workshop gives you the numbers; this structure gives you the narrative.

8. Common Mistakes in AI Cost-Benefit Analysis

Leaders who complete this exercise sometimes make errors that weaken their business case or lead to disappointment during implementation.

Mistake	Why It Happens	How to Avoid It
Underestimating current time spent	Normalisation — repetitive tasks become invisible when everyone has always done them	Ask team members to log actual time for two weeks; the results consistently exceed estimates
Using base salary instead of fully loaded cost	HR typically communicates salary, not total cost	Use the 1.4x multiplier or request the loaded figure from finance
Overestimating efficiency gains	Enthusiasm after seeing case studies with 70–85% improvements	Use 40% for initial projections; upgrade after pilot data confirms higher gains

Ignoring change management costs	Focus on tool cost overlooks the human adoption effort	Budget for training (8–16 hours per user), workflow redesign (project sponsor time), and a 30–60 day learning curve
Presenting savings without a specific action plan	The calculator produces impressive numbers that feel sufficient on their own	Always pair the financial analysis with a specific pilot proposal, timeline, and success criteria

9. Beyond the Calculator: Second-Order Benefits

The cost-saving calculator captures direct labour savings — the most easily quantifiable benefit. But several second-order benefits often prove equally or more valuable over time.

Quality and Consistency Improvement

Repetitive tasks performed by humans inevitably produce errors: transcription mistakes, inconsistent formatting, missed steps in multi-stage processes. AI performs these tasks identically every time. The value of this consistency shows up in reduced rework, fewer customer complaints, fewer compliance findings, and lower professional liability exposure. These benefits are real but harder to quantify in advance — they show up in your metrics 6–12 months after deployment.

Speed and Responsiveness

When your team spends 40% less time on repetitive tasks, your organisation becomes measurably faster. Proposals that took a week to generate take two days. Customer inquiries that waited in queue for hours get responses in minutes. Compliance reports that required end-of-quarter sprints can be produced on demand. Speed is a competitive advantage that does not appear directly in the cost-saving calculator but profoundly affects win rates, customer satisfaction, and market positioning.

Employee Satisfaction and Retention

Knowledge workers consistently report that repetitive tasks are the least satisfying part of their jobs. Removing 40% of this work does not just save money — it improves job satisfaction, reduces burnout, and increases retention. Given that replacing a knowledge worker costs 50–200% of their annual salary, retaining even one employee who would otherwise have left due to tedious workload more than pays for the AI investment.

10. The 30-Day Action Framework

The workshop concludes by asking you to complete the sentence: 'In the next thirty days, I will implement AI for...' This commitment is powerful, but only if backed by a concrete action plan.

Week	Action	Deliverable
Week 1	Identify your highest-cost repetitive task category and select 2–3 AI tools to evaluate	Shortlist of tools with pricing and feature comparison
Week 2	Run free trials or demos with your team on the actual tasks they perform daily	Team feedback on usability, accuracy, and time savings
Week 3	Build the business case using the calculator	One-page business case with ROI,

	framework and present to budget approver	payback period, and pilot plan
Week 4	Secure approval and begin pilot deployment on one specific process	Tool subscription active; 3–5 team members using it daily on designated tasks

The most important word in this framework is 'specific'. Vague commitments ('we will explore AI') produce no results. Specific commitments ('we will deploy Microsoft Copilot for our contract review team by July 15') create accountability and momentum.

11. Key Terms Glossary

Term	Definition
Fully Loaded Cost	The total cost of an employee to the organisation, including salary, benefits, payroll taxes, and overhead — typically 1.3x to 1.65x base salary
ROI (Return on Investment)	A measure of investment profitability, calculated as (net gain / cost of investment) x 100, expressed as a percentage
Payback Period	The time required for an investment to generate enough savings or revenue to recoup its cost — shorter is better
Blended Rate	A weighted average hourly cost across employees of different levels, used when a task involves multiple roles
Efficiency Gain	The percentage reduction in time or cost required to perform a task after implementing a new tool or process
Change Management	The structured approach to transitioning individuals, teams, and organisations from a current state to a desired future state — essential for AI adoption success
Pilot Deployment	A controlled, limited-scope implementation of a new tool or process, designed to test effectiveness and gather data before full-scale rollout
Second-Order Benefits	Indirect benefits that emerge over time from an initiative — such as improved retention, speed, or quality — beyond the directly measured financial savings

12. Quick Revision Checklist

- The cost-saving calculator quantifies AI's potential impact using your organisation's actual costs and task volumes
- Three task categories: document processing, routine communication, and data work — most organisations underestimate time spent on these by 30–40%
- Use fully loaded employee cost (salary x 1.4), not base salary, for accurate projections
- The savings formula: (Weekly hours x 52 x hourly rate x 0.40) minus annual AI tool cost = net savings
- Enterprise AI tools typically cost \$25–\$200 per user per month; \$80/month is a reasonable planning benchmark
- ROI of 500%+ and payback periods under 3 months are common for well-targeted AI deployments

- A strong business case includes: quantified problem, specific solution, financial impact, risk mitigation (pilot), and a clear ask
- Common mistakes: underestimating current time, using base salary, overestimating efficiency, ignoring change management
- Second-order benefits (quality, speed, retention) often match or exceed direct cost savings over 12+ months
- The 30-day action framework turns workshop insights into organisational results: shortlist tools, trial, build case, deploy pilot

20 Exam-Based MCQs — Cost-Saving Calculator

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. The cost-saving calculator workshop uses a 40% efficiency gain assumption. Why was this specific figure chosen?

- A) It is the maximum efficiency gain AI can deliver
- B) It is deliberately conservative — below the 60–85% gains documented in case studies — to build credibility with sceptical finance teams
- C) It is mandated by industry accounting standards for AI projections
- D) It represents the average across all AI deployments globally

Answer: B **Explanation:** B is correct — the 40% figure is intentionally conservative to ensure the business case is credible and defensible. Case studies like JPMorgan (87% time reduction), Klarna (700-employee equivalent), and Walmart consistently show higher gains. Using a lower figure builds trust with finance teams and creates room for upside surprise. A is wrong because documented gains frequently exceed 60%. C is fabricated. D is not the basis for the choice.

Q2. *Scenario:* A marketing director has a team of 15 people. She estimates they spend 10 hours per week collectively on creating routine social media posts and email newsletters. The average team member earns \$65,000 per year. She wants to calculate the annual cost of this repetitive communication work. Using the fully loaded cost method from the workshop, what is the approximate annual cost?

- A) \$16,900
- B) \$23,660
- C) \$33,800
- D) \$47,320

Answer: B **Explanation:** Fully loaded hourly rate = $(\$65,000 \times 1.4) / 2,000 = \$91,000 / 2,000 = \$45.50/\text{hour}$. Annual hours = 10 hours/week $\times$ 52 weeks = 520 hours. Annual cost = 520 $\times$ \$45.50 = \$23,660. A uses base salary without the 1.4 multiplier. C doubles the correct figure. D quadruples it.

Q3. Why does the workshop emphasise using fully loaded employee cost rather than base salary?

- A) To inflate the savings numbers and make AI look more attractive
- B) Because base salary alone understates the true cost by 25–40%, since it excludes benefits, taxes, and overhead that the organisation pays for each hour of work
- C) Because HR departments prefer to report fully loaded costs
- D) Because base salary information is confidential and harder to obtain

Answer: B **Explanation:** B is correct — the fully loaded cost captures every expense the organisation incurs for each employee hour (salary, benefits, payroll taxes, overhead). Using base salary alone systematically understates costs and produces artificially low ROI projections, which paradoxically makes it harder to get budget approval. A mischaracterises the purpose — accuracy, not inflation. C has no basis. D is wrong; salary data is readily available internally.

Q4. Scenario: A finance team of 8 people spends an estimated 25 hours per week collectively on data entry, reconciliation, and report generation. Their average fully loaded cost is \$55/hour. An AI tool costing \$60/user/month is being evaluated. The CFO wants to see the ROI. What is the approximate annual ROI of this AI investment?

- A) 250%
- B) 495%
- C) 890%
- D) 1,250%

Answer: B **Explanation:** Annual hours = $25 \times 52 = 1,300$. Annual cost = $1,300 \times \$55 = \$71,500$. Savings at 40% = $\$28,600$. Tool cost = $8 \times \$60 \times 12 = \$5,760$. Net savings = $\$28,600 - \$5,760 = \$22,840$. ROI = $(\$28,600 / \$5,760) \times 100 = 496.5\%$, approximately 495%. A is too low. C and D are too high and likely miscalculate the divisor.

Q5. Which of the three task categories typically consumes the MOST professional time in knowledge-work organisations?

- A) Routine communication
- B) Data work
- C) Document processing, which consumes 20–35% of total professional time according to McKinsey's workplace productivity analysis
- D) All three categories consume exactly equal amounts of time

Answer: C **Explanation:** C is correct — McKinsey's 2023 analysis found document processing (reading, reviewing, creating, extracting data from documents) consumes 20–35% of total professional time in knowledge-work organisations, making it typically the largest single category. A and B also consume significant time but generally less than document processing at the organisational level. D is incorrect — the distribution varies by function and organisation.

Q6. Scenario: After completing the cost-saving calculator, a VP of Operations identifies \$180,000 in potential annual savings from AI automation of data entry and report generation. She presents the

business case to her CFO, who responds: 'These numbers look too good. I need proof before I approve \$30,000 for AI tools.' What is the MOST effective response?

- A) Show the CFO case studies from other companies to prove that the savings are realistic
- B) Propose a 60-day pilot on one specific process with measurable success criteria, requiring only a fraction of the \$30,000 commitment
- C) Accept the CFO's rejection and abandon the initiative
- D) Go ahead with implementation without CFO approval since the ROI is clearly positive

Answer: B **Explanation:** B is correct because a limited pilot addresses the CFO's core concern (proof) while minimising risk. A free trial or single-month subscription for a pilot costs little and produces the real-world evidence the CFO wants. This is the risk mitigation step recommended in the business case framework. A provides external evidence but not proof from the organisation's own operations, which is what the CFO is asking for. C abandons a strong opportunity unnecessarily. D violates basic organisational governance.

Q7. The workshop identifies that most leaders underestimate their team's time on repetitive tasks by 30–40%. What is the RECOMMENDED way to get an accurate figure?

- A) Trust the leader's initial estimate since they know their team best
- B) Ask team members to log their actual time on repetitive tasks for two weeks
- C) Use industry benchmarks without any internal validation
- D) Hire a consulting firm to conduct a time-and-motion study

Answer: B **Explanation:** B is correct — actual time logging over two weeks consistently reveals higher numbers than estimates because normalised, daily repetitive work becomes invisible to leaders. The results provide credible, organisation-specific data for the business case. A produces underestimates due to normalisation bias. C ignores organisation-specific realities. D is disproportionate and expensive for this purpose.

Q8. *Scenario:* A company deploys AI for routine customer service communication. After three months, direct time savings are \$45,000. However, the customer satisfaction (CSAT) score has improved from 72 to 84, employee turnover in the customer service team has dropped from 35% to 18%, and average response time has decreased from 4 hours to 22 minutes. These improvements beyond direct time savings are BEST described as:

- A) Coincidences unrelated to the AI deployment
- B) Second-order benefits — indirect improvements in quality, speed, and retention that emerge over time and often match or exceed direct cost savings
- C) Marketing metrics that should not be included in ROI calculations
- D) Signs that the AI deployment was poorly scoped and delivered unintended results

Answer: B **Explanation:** B is correct — second-order benefits include quality improvements (higher CSAT), speed gains (faster response times), and employee satisfaction (lower turnover). These benefits often match or exceed direct cost savings over 12+ months. They are real, measurable, and should be included in the full value assessment. A ignores the causal connection between AI automation and these

improvements. C artificially excludes material benefits. D mischaracterises positive outcomes as problems.

Q9. In the 30-day action framework, what should be completed during Week 3?

- A) Full-scale deployment of the AI tool across the entire organisation
- B) Building the business case using the calculator framework and presenting it to the budget approver
- C) Identifying the highest-cost repetitive task category
- D) Running free trials with the team

Answer: B **Explanation:** B is correct — Week 3 is dedicated to building and presenting the business case with ROI, payback period, and pilot plan. C is a Week 1 activity (identify tasks and shortlist tools). D is a Week 2 activity (trials and team feedback). A is premature; Week 4 begins a pilot, not full deployment.

Q10. *Scenario:* An HR director calculates that her team of 12 spends 30 hours per week on resume screening, interview scheduling, and candidate communication. Average salary is \$60,000. She is excited about an AI tool costing \$100/user/month but realises she forgot to account for change management costs in her business case. Which change management costs should she include for a realistic projection?

- A) No change management costs are necessary — AI tools are intuitive and require no training
- B) Training time (8–16 hours per user), workflow redesign effort (project sponsor time), and a 30–60 day learning curve during which savings will be lower than projected
- C) Only the cost of a one-hour webinar to introduce the tool
- D) The cost of hiring a full-time change management consultant for 12 months

Answer: B **Explanation:** B is correct — realistic change management budgeting includes training time (each team member needs 8–16 hours to learn the tool), workflow redesign (someone needs to map how the tool integrates into existing processes), and an acknowledged ramp-up period where efficiency gains are lower than the steady-state projection. A is naive — even intuitive tools require workflow adaptation. C dramatically understates the effort. D is disproportionate for this scale of deployment.

Q11. Using the workshop formula, what is the payback period for a \$24,000 annual AI tool investment that generates \$8,000 per month in savings?

- A) 3 months
- B) 6 months
- C) 12 months
- D) 24 months

Answer: A **Explanation:** Payback period = Annual tool cost / Monthly savings = \$24,000 / \$8,000 = 3 months. The full annual cost is recouped in three months; the remaining nine months generate pure net savings. B, C, and D overstate the payback period.

Q12. Scenario: Two department heads present AI business cases to the CFO on the same day. Department A shows a \$500,000 savings opportunity with a 200% ROI. Department B shows a \$50,000 savings opportunity with a 1,200% ROI. The CFO has budget for only one pilot. Which pilot should the CFO prioritise, and why?

- A) Department A, because \$500,000 is a larger absolute savings
- B) Department B, because 1,200% ROI indicates dramatically more efficient use of investment dollars
- C) It depends on additional factors: the confidence level of each projection, the strategic importance of each function, the ease of implementation, and the organisation's risk appetite
- D) Neither — the CFO should reject both proposals until market conditions improve

Answer: C **Explanation:** C is correct — ROI and absolute savings alone do not tell the full story. A \$500K savings with 200% ROI requires a \$250K investment but carries implementation complexity. A \$50K savings with 1,200% ROI requires only ~\$4K investment but has smaller total impact. The right choice depends on which projection is more reliable, which function is more strategically important, and which team is more ready for AI adoption. A and B each fixate on one metric. D is unnecessarily risk-averse.

Q13. The Deloitte 2023 AI ROI Survey found that organisations with a clear, quantified business case were how much more likely to achieve their target ROI?

- A) 10% more likely
- B) 50% more likely
- C) 2.4 times more likely
- D) 5 times more likely

Answer: C **Explanation:** C is correct — the Deloitte survey found that organisations which started with a clear, quantified business case were 2.4 times more likely to achieve their target ROI than those that deployed AI without specific financial targets. This underscores the value of the calculator exercise. A and B understate the difference. D overstates it.

Q14. Scenario: A logistics company has completed the cost-saving calculator and identified three opportunity areas: document processing (\$120,000 annual savings potential), route communication (\$45,000), and shipment data entry (\$85,000). The CEO asks which area to target first. What should drive the prioritisation decision?

- A) Always start with the highest-savings category regardless of other factors
- B) Start with whichever category has the most enthusiastic team, since adoption is the primary success factor
- C) Start with the highest-cost category, as recommended in the workshop — this is where you will see the fastest ROI and most compelling proof of concept
- D) Start with all three simultaneously to maximise total savings in year one

Answer: C **Explanation:** C is correct — the workshop specifically recommends starting with the highest-cost category because it produces the most visible results and the strongest proof point for expansion. Document processing at \$120K annual potential is the clear priority. A reaches the same conclusion but for

a less nuanced reason. B is a valid secondary consideration but not the primary driver. D spreads resources too thin and complicates measurement — a common pitfall the workshop warns against.

Q15. Why does the cost-saving calculator deliberately exclude second-order benefits from its core calculation?

- A) Because second-order benefits are not real and should not be counted
- B) Because they are harder to quantify in advance and including uncertain projections weakens the credibility of the business case with finance teams
- C) Because accounting standards prohibit including indirect benefits in investment analyses
- D) Because second-order benefits only appear after 5+ years

Answer: B **Explanation:** B is correct — second-order benefits (quality improvement, speed gains, retention improvement) are real but harder to predict with precision before deployment. Including uncertain numbers weakens the credibility of an otherwise rigorous business case. The strategy is to present conservative, defensible direct savings and let second-order benefits serve as upside when they materialise 6–12 months later. A is wrong — they are real. C is fabricated. D is wrong — they typically emerge within 6–12 months.

Q16. What is the MOST important word in the 30-day action framework commitment sentence 'In the next thirty days, I will implement AI for...'?

- A) 'AI' — because it specifies the technology
- B) 'thirty days' — because it creates a time-bound commitment
- C) 'specific' (implied) — because vague commitments like 'explore AI' produce no results, while specific commitments like 'deploy Copilot for contract review by July 15' create accountability
- D) 'implement' — because it distinguishes action from research

Answer: C **Explanation:** C is correct — the workshop emphasises that specificity is the critical differentiator between commitments that produce results and those that do not. 'We will explore AI' is not a commitment; 'We will deploy Microsoft Copilot for our contract review team by July 15' is. Specificity creates accountability, measurability, and momentum. B and D are important elements but secondary to the specificity that makes the commitment actionable.

Q17. *Scenario:* A VP completes the calculator and finds a 600% ROI for deploying AI in document processing. She is about to present to the board but realises she has not accounted for the 30–60 day learning curve during which savings will be lower than projected. How should she adjust her business case?

- A) Remove the learning curve from her presentation since it complicates the narrative
- B) Acknowledge the ramp-up period, reduce Year 1 savings by 10–15% to account for the learning curve, and note that steady-state ROI will be achieved from month 3–4 onward
- C) Delay the presentation until after a pilot has been completed to show actual ramp-up data
- D) Double the learning curve estimate to be extra conservative

Answer: B **Explanation:** B is correct — acknowledging the learning curve while quantifying its impact (10–15% reduction in Year 1) shows rigour and builds credibility. The adjustment is material enough to mention but not large enough to change the investment decision. A is dishonest and will undermine credibility if a board member asks about adoption challenges. C delays action unnecessarily. D is overly conservative and may kill a good project.

Q18. The workshop states that replacing a knowledge worker costs 50–200% of their annual salary. Why is this figure relevant to the AI cost-saving calculator discussion?

- A) It justifies paying AI tools more than employees
- B) It shows that retaining even one employee who might have left due to tedious workload — by removing that tedium with AI — can more than pay for the entire AI tool investment
- C) It proves that AI will replace all knowledge workers within five years
- D) It is used to calculate the fully loaded employee cost

Answer: B **Explanation:** B is correct — if AI removes repetitive tasks that contribute to burnout and turnover, and this prevents even one resignation, the avoided replacement cost (\$35,000–\$140,000 for a \$70,000 employee) can exceed the entire annual AI tool investment. This is a second-order benefit that strengthens the overall value proposition. A makes no logical connection. C is wrong — the point is retention improvement, not replacement. D confuses two different concepts.

Q19. *Scenario:* After completing the workshop, a department head realises his initial estimate of 15 hours per week on repetitive tasks was far too low. When his team logged their actual time for two weeks, the real figure was 28 hours per week — 87% higher than his estimate. This discrepancy is BEST explained by:

- A) The team deliberately inflated their time logs to justify AI tool purchases
- B) Normalisation bias — when repetitive tasks are part of daily routine, leaders lose visibility into how much time they actually consume
- C) The team was working inefficiently and needs performance management, not AI tools
- D) The two-week logging period was unrepresentative of normal workload

Answer: B **Explanation:** B is correct — normalisation causes repetitive tasks to become invisible when they are embedded in daily routines. Leaders who have 'always' seen their teams doing these tasks stop perceiving the time as waste. This 30–40% underestimation is well-documented and is exactly why the workshop recommends actual time logging over estimates. A assumes bad faith without evidence. C assumes the tasks are unnecessary rather than that they are repetitive and automatable. D dismisses the data without justification.

Q20. A complete AI business case, according to the workshop framework, includes which five elements?

- A) Executive summary, market analysis, competitor benchmarking, financial projections, and appendix

- B) Problem quantified, solution proposed, financial impact calculated, risk mitigation plan (pilot), and a clear ask (budget and resources needed)
- C) Technology overview, vendor comparison, implementation timeline, training plan, and success metrics
- D) Vision statement, strategic alignment, resource requirements, governance framework, and change management plan

Answer: B **Explanation:** B is correct — the workshop's business case framework consists of: (1) the problem quantified with specific costs and hours, (2) the solution proposed with specific tools and scope, (3) the financial impact with ROI and payback period, (4) risk mitigation through a controlled pilot with success criteria, and (5) a clear ask specifying the budget and resources needed. A describes a general business plan, not a tool-specific business case. C covers implementation planning but misses the financial case. D is a strategic planning framework, not an investment justification.