

INTERACTIVE WORKSHOP: Your 90-Day Plan

Generative AI for Business Leaders & C-Suite | AI Implementation Planning

1. Why 90 Days Is the Right Time Horizon

Most AI initiatives fail not because the technology is wrong, but because the planning horizon is. Annual roadmaps create a false sense of security — leaders commission a year-long strategy document, circulate it for weeks of review, and by the time everyone signs off, the competitive window has closed. At the other extreme, ad-hoc experimentation without structure produces scattered pilots that never scale. Ninety days strikes the balance: long enough to move through education, piloting, and early scaling, but short enough to maintain urgency and accountability.

Research from McKinsey’s 2024 State of AI report found that organisations that moved from concept to live pilot within 90 days were 2.4 times more likely to report positive ROI than those with longer timelines. The 90-day arc also aligns with a single business quarter, which makes budget approval simpler and gives you a natural reporting cycle for executives and boards.

Key Point: The 90-day plan is not a compromise between speed and thoroughness — it is the minimum viable time to prove AI value while the maximum allowable time before strategic drift sets in.

2. The Four Phases in Detail

Your 90-day plan divides into four distinct phases, each with its own deliverables and success criteria. Understanding what belongs in each phase prevents the most common mistake: jumping straight to tool selection before the organisation is ready.

Phase	Timeframe	Primary Activities	Key Deliverable
1. Education & Exploration	Weeks 1–2	Executive alignment, AI literacy workshops, departmental pain-point mapping	Shortlist of 3–5 candidate pilot projects
2. Pilot Projects	Weeks 3–6	Tool selection, pilot scoping, implementation sprints, success-metric tracking	Documented pilot results with ROI data
3. Scaling Winners	Weeks 7–10	Expand proven pilots, cross-department rollout planning, budget requests	Scaling proposal with projected enterprise impact
4. Commitments & Accountability	Weeks 11–12	Personal action items, stakeholder communication, governance check	Signed-off 12-month roadmap draft

Phase 1: Education and Exploration (Weeks 1–2)

The first two weeks are about closing the knowledge gap and building organisational buy-in. You need an executive sponsor — someone whose authority can clear budget and political obstacles. If you are the C-suite sponsor, your job is to visibly champion the initiative rather than delegate it quietly to IT. Identify

three to five departments that will participate in initial exploration. Choose departments with measurable pain points and a willingness to experiment. The leadership AI workshop during this phase is non-negotiable: it gives every stakeholder a shared vocabulary and aligns expectations before any money is spent.

A practical tip: schedule the workshop within the first two weeks, not ‘sometime soon.’ Putting a specific date on the calendar converts good intentions into commitments. Organisations that schedule their workshop in week one launch pilots 40% faster than those that leave it open-ended, according to internal tracking data from Deloitte’s AI advisory practice.

Phase 2: Pilot Projects (Weeks 3–6)

Each pilot needs four elements defined up front: the specific task being automated (not ‘customer service’ but ‘responding to order-status inquiries’), a measurable success metric with a number attached, a named pilot owner with authority and time, and a specific AI tool. The tool choice matters less than you think at this stage — most enterprise AI platforms can handle common use cases — but making a decision forces clarity. You can always switch tools; you cannot recover the weeks lost to analysis paralysis.

Limit yourself to three to five pilots. More than five dilutes attention and increases the risk of none succeeding. Fewer than three leaves you without a portfolio effect: if one pilot hits an obstacle, you still have two others generating results to sustain organisational momentum.

Phase 3: Scaling Winners (Weeks 7–10)

Not every pilot will succeed, and that is by design. By week seven you need explicit decision criteria: what threshold of improvement qualifies a pilot for scaling? A common benchmark is at least 30% efficiency improvement, user satisfaction above 8/10, and technical stability with minimal support overhead. Scaling also requires communication — lunch-and-learns, internal case-study emails, short demo videos — so the broader organisation sees evidence, not just promises.

Budget planning belongs in this phase, not later. If your pilots prove successful, requesting funding retroactively creates delays. Pre-approved contingency budgets (even modest ones) let you move immediately when a pilot earns its expansion.

Phase 4: Commitments and Accountability (Weeks 11–12)

The final phase converts your plan from a document into organisational behaviour. You need personal, dated commitments: who will you share the plan with (names, not roles), what budget approval is needed and from whom, what is the single biggest obstacle and your strategy to overcome it, and what will you do literally tomorrow. Public commitment increases follow-through dramatically. Research by Gail Matthews at Dominican University found that people who write down goals and share them with an accountability partner achieve 33% more of their goals than those who merely think about them.

3. The Executive Sponsor Role

The executive sponsor is the single most predictive factor in whether an AI initiative reaches scale. A 2023 Boston Consulting Group study of 250 enterprise AI projects found that initiatives with an active C-level sponsor were 3.5 times more likely to scale beyond pilot phase. ‘Active’ is the operative word — a

name on a slide deck is not sponsorship. Active sponsorship means attending governance meetings, removing bureaucratic obstacles, publicly communicating the initiative’s importance, and protecting the team when early results are mixed.

If you are the executive sponsor, resist the temptation to over-delegate. Your visible involvement signals to the organisation that AI is a strategic priority, not a side experiment. If you need to recruit a sponsor above you, prepare a one-page brief: the problem, the proposed pilot, the expected ROI, and the ask (time, budget, air cover). Executives respond to concise, quantified proposals.

Real-World Incident — Moderna (2023): Moderna’s CEO Stéphane Bancel personally championed the company’s AI rollout, setting a public goal of integrating AI into every business function within 90 days. His visible sponsorship led to over 750 GPT-based applications being deployed company-wide, touching drug discovery, manufacturing, legal, and HR. Without C-suite sponsorship at that level, middle-management caution would have stalled the initiative at the pilot stage.

4. Selecting and Scoping Pilot Projects

The art of pilot selection is choosing projects that are valuable enough to matter but contained enough to succeed quickly. A useful mental model is the Impact–Feasibility Matrix:

	High Feasibility	Low Feasibility
High Impact	Start here — ideal pilots	Strategic — plan for Phase 2 or 3
Low Impact	Quick wins — use to build confidence	Avoid — high effort, low return

Good pilot characteristics: the task is repetitive and time-consuming, a clear before/after metric exists (time saved, error rate, cost), the data needed is already available (no six-month data-cleaning project first), and a willing department owner is identified. Bad pilot characteristics: the task requires complex judgment with no clear right answer, success depends on integrating five systems simultaneously, or the department is resistant and will passively sabotage the experiment.

Writing Success Metrics That Work

Vague metrics like ‘improve efficiency’ are useless. Every pilot needs a number and a deadline. ‘Reduce average response time for order-status inquiries from 4 hours to under 30 minutes within 4 weeks’ is a metric. ‘Make customer service better’ is a wish. The metric should be something the pilot owner can measure without building new reporting infrastructure — if measuring success requires a separate data-engineering project, you have scoped the pilot too broadly.

5. Tool Selection: Build vs. Buy in Practice

For 90-day plans, the answer is almost always ‘buy.’ Building custom AI solutions takes months of development and data science resources that most organisations do not have. Off-the-shelf enterprise AI tools — ChatGPT Enterprise, Microsoft Copilot, Google Gemini for Workspace, Anthropic’s Claude for Business, Intercom’s AI agent, Gong’s revenue intelligence — cover the vast majority of initial use cases. The decision matrix below helps you choose:

Use Case Category	Recommended Tool Type	Typical Cost Range
General writing, summarisation, analysis	ChatGPT Enterprise, Claude, Copilot	\$20–\$60/user/month
Customer service automation	Intercom, Zendesk AI, Ada	\$50–\$200/agent/month
Sales intelligence & coaching	Gong, Clari, Salesforce Einstein	\$100–\$200/user/month
Code generation & developer productivity	GitHub Copilot, Cursor, Amazon CodeWhisperer	\$10–\$40/user/month
Marketing content & creative	Jasper, Writer, Adobe Firefly	\$40–\$150/user/month
Data analysis & BI	Tableau AI, Power BI Copilot, ThoughtSpot	Varies by licence

Make a tool decision within the first two weeks. Vendor evaluation paralysis is one of the top three reasons 90-day plans stall. Pick the leading option in your use-case category, negotiate a short-term pilot licence (most vendors offer 30-day trials for enterprise), and move to implementation. You can switch tools later; you cannot recover lost time.

6. Budget Planning and Stakeholder Buy-In

AI pilots are inexpensive relative to other enterprise technology investments, but they still require explicit budget. A typical three-pilot 90-day plan might include: enterprise AI tool licences for 20–50 users (\$2,000–\$8,000 for three months), a half-day leadership workshop facilitated externally (\$3,000–\$8,000), implementation support or consulting hours (\$5,000–\$15,000), and internal time allocation for the pilot owner and participants (opportunity cost, not cash). Total cash outlay for a meaningful pilot programme: \$10,000–\$30,000 — a rounding error in most enterprise budgets, yet enough to produce real, measurable results.

Frame the budget conversation around risk, not cost. The question is not ‘Can we afford \$20,000 for an AI pilot?’ but ‘Can we afford to be the last company in our sector without AI capabilities?’ Connect the budget request to a specific competitive threat or revenue opportunity identified in earlier sections of the course.

7. Overcoming the Top 5 Implementation Obstacles

Every 90-day plan encounters friction. Anticipating obstacles and pre-loading solutions is what separates execution from aspiration.

Obstacle	Root Cause	Pre-Loaded Solution
Sceptical colleagues	Fear of job displacement or change fatigue	Frame AI as augmenting their role; show early wins that make their lives easier, not harder
Limited IT resources	IT backlog and competing priorities	Use SaaS AI tools that require minimal IT involvement; cloud-hosted, no infrastructure
Competing priorities	AI is ‘important but not	Tie AI to a measurable quarterly goal the executive

	urgent'	team already cares about
Data readiness concerns	'Our data isn't clean enough' becomes an excuse for inaction	Start with use cases that work on unstructured text (emails, documents) rather than requiring clean databases
Governance uncertainty	'Legal hasn't approved it yet'	Draft lightweight policies in week one (covered in Block 5); perfection is the enemy of progress

8. The Accountability Mechanism

Writing a plan and executing a plan are separated by one variable: accountability. The workshop builds in three accountability layers. First, personal commitments with specific next-day actions — not 'next week' or 'soon,' but literally tomorrow. Immediate action creates momentum. Second, sharing the plan with at least one named person who will follow up. A plan that lives only in your notebook has a survival rate close to zero. Third, connecting the plan to a natural quarterly reporting cycle so progress is reviewed in a forum that already exists.

The accountability partner does not need to be a supervisor. A peer, a mentor, or a colleague going through the same exercise works well. The point is external visibility: when someone else knows your commitment, the social cost of inaction rises.

Key Point: Public commitment with a specific timeline increases goal achievement by approximately one-third compared to private intentions. Share your plan within 48 hours of completing it.

9. From Plan to Roadmap: Connecting 90 Days to 12 Months

The 90-day plan is not the destination — it is the on-ramp. By the end of week 12, you should have enough evidence to draft a 12-month AI roadmap that includes: which pilots to scale enterprise-wide and on what timeline, what new use cases to pilot next (informed by what you learned), the governance framework refinements needed for broader deployment, the training and hiring investments required, and a projected ROI model grounded in actual pilot data rather than vendor marketing.

The transition from 90-day plan to 12-month roadmap is also where you shift from 'proving AI works' to 'embedding AI in strategy.' The 12-month roadmap should be reviewed by the board or executive committee and become a standing agenda item in quarterly business reviews.

10. Workshop Template Walkthrough

The workshop template used in the lecture contains four sections. Here is an expanded guide to filling each one with rigour:

Section 1: Education & Exploration (Weeks 1–2)

- **Executive sponsor:** Write a specific name and title. If it is you, state your authority level. If you need to recruit one, note the specific meeting where you will ask.

- **Participating departments:** List at least three. Prioritise departments with quantifiable pain points (e.g., customer service handling 500 tickets/day manually) over departments that are merely ‘interested.’
- **Workshop date:** A calendar date within 14 days, not a week number. Block the room and send invitations before moving to the next section.
- **Candidate pilot projects:** Three to five ideas. For each, note the department, the specific process, and a rough success metric. These are hypotheses, not commitments.

Section 2: Pilot Projects (Weeks 3–6)

- **Specific task:** Narrow and precise. ‘Drafting first responses to Tier 1 support tickets’ not ‘improving customer service.’
- **Success metric:** A number with a unit and a deadline. ‘Reduce response time from 4 hours to 30 minutes within 4 weeks.’
- **Pilot owner:** A named individual with time, authority, and enthusiasm. This person will report weekly on progress.
- **AI tool:** A specific product, not ‘some AI tool.’ Make the best available decision now; you can change it.

Section 3: Scaling Winners (Weeks 7–10)

- **Decision criteria:** Write numeric thresholds. E.g., $\geq 30\%$ efficiency gain, user satisfaction $\geq 8/10$, support burden < 2 hours/week.
- **Communication plan:** Choose channels: all-hands demo, internal newsletter case study, Slack showcase channel, video testimonial from pilot owner.
- **Scaling budget:** Estimate licence costs for expanded user base, training for new departments, and a part-time implementation coordinator.
- **Priority departments for scaling:** Sequence matters. Scale within the pilot department first (deep), or spread across departments (wide)? Decide based on where the biggest ROI lies.

Section 4: Personal Commitments

- **Share with three people:** Your manager or board, your leadership team, and one external peer or mentor. Meetings should be calendared within seven days.
- **Budget approval:** State the exact amount, the approver’s name, and the meeting where you will present the request.
- **Biggest obstacle and strategy:** Be honest. Name the obstacle and write a one-sentence mitigation. If the obstacle is political, your strategy should involve people, not technology.
- **Tomorrow’s action:** One specific, completable action. ‘Send the email to the VP of Customer Service proposing a 15-minute meeting about the AI pilot.’

11. Key Terms Glossary

Term	Definition
90-Day AI Plan	A structured, time-bound implementation roadmap covering education, piloting, scaling, and accountability across a single business quarter

Executive Sponsor	A C-level or senior leader who visibly champions and resource-protects an AI initiative
Pilot Project	A small-scale, controlled implementation of an AI solution designed to test feasibility and measure value before scaling
Success Metric	A specific, quantified, time-bound measure used to evaluate whether a pilot achieved its intended outcome
Impact–Feasibility Matrix	A 2×2 framework for prioritising projects by potential business value (impact) and ease of implementation (feasibility)
Scaling Criteria	Pre-defined numeric thresholds that a pilot must meet before it is approved for broader organisational deployment
Accountability Partner	A person who has visibility into your commitments and will follow up on progress, increasing the probability of execution
Enterprise AI Tool	A commercial AI platform with contractual data-protection guarantees, suitable for handling business-confidential information
Quick Win	A high-feasibility, lower-impact project used to build organisational confidence and demonstrate AI capability early
Governance Framework	The policies, structures, and processes that ensure AI is used safely, ethically, and in compliance with regulations

12. Quick Revision Checklist

- The 90-day plan is divided into four phases: Education & Exploration, Pilot Projects, Scaling Winners, and Commitments & Accountability.
- An active executive sponsor is the strongest predictor of whether an AI initiative scales beyond pilot phase.
- Limit pilots to three to five projects, each with a named owner, a specific task, a numeric success metric, and a chosen tool.
- Use the Impact–Feasibility Matrix to select pilots: start with high-impact, high-feasibility candidates.
- Success metrics must be numeric and time-bound — ‘reduce X by Y% within Z weeks.’
- Budget for a 90-day AI pilot programme typically ranges from \$10,000 to \$30,000 — position it as risk mitigation, not expense.
- Pre-load solutions for the five most common obstacles: sceptical colleagues, limited IT resources, competing priorities, data readiness, and governance uncertainty.
- Accountability increases goal achievement by approximately one-third — share your plan with a named person within 48 hours.
- The 90-day plan feeds into a 12-month roadmap grounded in real pilot data, not projections.
- Schedule the leadership workshop within two weeks and make a tool decision within the same timeframe — speed matters more than perfection at this stage.

20 Exam-Based MCQs — Your 90-Day AI Implementation Plan

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. Scenario: A VP of Operations has just completed a leadership AI workshop and wants to move quickly. She asks her team to identify 12 potential AI pilot projects across all departments so they can 'cover maximum ground.' Her CTO advises limiting the initial portfolio to three to five projects. Why is the CTO's advice strategically sound?

- A) Fewer pilots reduce software licensing costs below the approval threshold
- B) More than five pilots dilutes attention and increases the risk that none succeeds, eliminating the portfolio effect
- C) Enterprise AI tools cannot support more than five concurrent use cases technically
- D) Regulatory frameworks in most industries cap AI deployments at five per quarter

Answer: B **Explanation:** B is correct because spreading resources across too many pilots means none gets enough focus to produce compelling results, destroying organisational momentum. A is wrong because cost is not the primary concern — the issue is execution focus. C is wrong because enterprise tools have no such technical limitation. D is wrong because no such regulatory cap exists.

Q2. In the 90-day AI implementation framework, what is the PRIMARY purpose of Weeks 1–2?

- A) Selecting and licensing AI tools for the organisation
- B) Building executive alignment, closing knowledge gaps, and identifying candidate pilot projects
- C) Deploying AI solutions to customer-facing departments
- D) Drafting a 12-month AI roadmap for board approval

Answer: B **Explanation:** B is correct because the first two weeks focus on education, executive sponsorship, and identifying which departments and projects to pilot. A is wrong because tool selection happens in Phase 2 alongside pilot scoping. C is wrong because deployment occurs during the pilot phase (weeks 3–6). D is wrong because the 12-month roadmap is a post-90-day output informed by pilot results.

Q3. Scenario: A Director of Customer Experience defines her pilot's success metric as 'improve customer satisfaction.' Her AI consultant recommends revising this before launching the pilot. What is the MOST important improvement needed in this metric?

- A) The metric should reference a specific AI tool to ensure vendor accountability
- B) The metric should include a numeric target and a deadline, such as 'increase CSAT from 72% to 85% within 4 weeks'
- C) The metric should focus on cost reduction rather than customer satisfaction
- D) The metric should be approved by the legal department before the pilot begins

Answer: B **Explanation:** B is correct because a success metric without a number and a deadline is unmeasurable — you cannot determine whether the pilot succeeded or failed. A is wrong because the metric measures outcomes, not tool usage. C is wrong because customer satisfaction is a valid metric; the problem is vagueness, not the category. D is wrong because legal approval of a metric is not standard practice.

Q4. Scenario: A mid-size insurance company's AI initiative has the CEO's verbal support but no active sponsorship — the CEO has not attended any governance meetings, cleared budget obstacles, or communicated the initiative's importance to department heads. The pilot team is struggling to get resources from IT and cooperation from the claims department. What should the project lead do FIRST?

- A) Proceed without executive involvement and demonstrate results to earn support later
- B) Prepare a one-page brief for the CEO quantifying the ask and request a 15-minute meeting to secure active sponsorship
- C) Escalate to the board of directors to override the CEO's passive stance
- D) Pivot the pilot to a department that does not require IT support

Answer: B **Explanation:** B is correct because the core problem is passive sponsorship — converting verbal support into visible, active backing requires a concise, quantified request that respects the CEO's time. A is wrong because without executive air cover, resource and political obstacles will persist. C is wrong because going over the CEO's head is politically destructive and premature. D is wrong because it treats a symptom (IT resistance) rather than the root cause (lack of executive sponsorship).

Q5. Which characteristic makes a task a POOR candidate for a 90-day AI pilot project?

- A) The task is repetitive, time-consuming, and has a clear before/after metric
- B) The task requires complex judgment with no clear right answer and depends on integrating five systems simultaneously
- C) The data needed for the task is already available in existing systems
- D) A willing department owner has been identified to lead the pilot

Answer: B **Explanation:** B is correct because tasks requiring complex, subjective judgment and extensive system integration are high-risk, low-feasibility projects that are likely to stall within 90 days. A, C, and D all describe characteristics of good pilot candidates: repeatability, data readiness, and an engaged owner.

Q6. Scenario: A pharmaceutical company's head of R&D proposes building a custom large language model trained on the company's proprietary clinical trial data as the first AI initiative. The timeline estimate is six months of data preparation and four months of model training. The CTO suggests using an off-the-shelf enterprise AI tool instead. In the context of a 90-day plan, which approach is MORE appropriate and why?

- A) The custom model, because proprietary data creates a competitive moat that off-the-shelf tools cannot replicate
- B) The off-the-shelf tool, because it can deliver measurable results within the 90-day window, while the custom model's 10-month timeline violates the plan's urgency principle
- C) Neither — the company should wait for regulatory clarity on AI in pharmaceuticals before starting any initiative
- D) The custom model, because enterprise AI tools are not suitable for regulated industries

Answer: B **Explanation:** B is correct because the 90-day framework prioritises speed to measurable results; a 10-month build timeline defeats the purpose. The custom model may be strategically valid later, but it is not a starter pilot. A is wrong because the competitive moat argument, while valid long-term, does

not justify delaying all AI momentum. C is wrong because regulatory clarity is not a prerequisite for internal experimentation with approved tools. D is wrong because enterprise AI tools are widely used in regulated industries under appropriate governance.

Q7. What is the recommended total cash outlay for a typical three-pilot, 90-day AI implementation programme?

- A) Under \$1,000
- B) \$10,000–\$30,000
- C) \$200,000–\$500,000
- D) \$1 million or more

Answer: B **Explanation:** B is correct because a typical programme includes enterprise tool licences (\$2,000–\$8,000), a workshop (\$3,000–\$8,000), and implementation support (\$5,000–\$15,000), totalling \$10,000–\$30,000. A is too low to cover even basic tool licensing. C and D reflect enterprise-scale custom AI development, not a 90-day pilot programme.

Q8. *Scenario:* A retail chain's AI pilot reduced customer email response time from 6 hours to 20 minutes, with a user satisfaction score of 9.1/10 among customer service agents. The pilot ran for four weeks in one region covering 800 tickets. The VP of Operations wants to know if this is ready to scale. Based on the scaling criteria discussed in the framework, is this pilot ready to scale?

- A) No, because a four-week pilot is too short to draw scaling conclusions
- B) No, because user satisfaction should be measured among customers, not agents
- C) Yes, because it exceeds typical thresholds: >30% efficiency improvement and user satisfaction above 8/10
- D) Yes, but only if the vendor provides a written guarantee of identical performance at scale

Answer: C **Explanation:** C is correct because the pilot exceeds the standard scaling benchmarks (efficiency improved by over 90%, and agent satisfaction is 9.1/10). A is wrong because four weeks is within the pilot phase timeframe and sufficient for repeatable tasks. B is wrong because agent satisfaction is a valid metric; customer satisfaction can be added during scaling. D is wrong because vendor guarantees of identical performance at scale are unrealistic and unnecessary if the pilot data is strong.

Q9. According to research cited in the course, by approximately how much does sharing goals with an accountability partner increase goal achievement?

- A) 10%
- B) 33%
- C) 75%
- D) It has no statistically significant effect

Answer: B **Explanation:** B is correct, referencing Gail Matthews' research at Dominican University showing a 33% increase in goal achievement when goals are written down and shared. A understates the effect. C overstates it. D is contradicted by the research.

Q10. Scenario: A financial services firm's AI steering committee is evaluating pilot proposals. The compliance department wants to automate regulatory report generation (high impact, high feasibility). The innovation lab wants to build an AI-powered trading algorithm (high impact, low feasibility). The HR department wants AI to screen job postings for inclusive language (low impact, high feasibility). Using the Impact–Feasibility Matrix, which project should be the FIRST pilot?

- A) The AI-powered trading algorithm, because it has the highest potential revenue impact
- B) The inclusive language screener, because it has the lowest risk
- C) The regulatory report generator, because it sits in the high-impact, high-feasibility quadrant
- D) All three should be launched simultaneously to maximise learning

Answer: C **Explanation:** C is correct because the high-impact, high-feasibility quadrant represents ideal starter pilots. A is wrong because low feasibility means it is unlikely to produce results within 90 days. B is wrong because low-impact projects, while useful as quick wins, should not be the lead pilot when a high-impact option is available. D is wrong because launching all three dilutes focus, especially given that the trading algorithm is low-feasibility.

Q11. Why should budget for scaling be estimated during the pilot phase (Weeks 7–10) rather than after pilots conclude?

- A) Because AI tool vendors require advance payment for enterprise licences
- B) Because pre-approved contingency budgets eliminate the delay between proving a pilot works and expanding it
- C) Because finance departments do not process budget requests outside of quarterly planning cycles
- D) Because scaling costs are easier to estimate before pilot data is available

Answer: B **Explanation:** B is correct because the window of organisational enthusiasm after a successful pilot is narrow; waiting for budget approval during that window wastes momentum. A is wrong because most vendors offer flexible payment terms. C is wrong because finance departments routinely handle out-of-cycle requests for strategic initiatives. D is wrong and backwards — pilot data makes estimation more accurate, not less.

Q12. Scenario: A manufacturing company's IT director pushes back on the AI pilot timeline, saying: 'Our data isn't clean enough for AI. We need at least six months to build a proper data warehouse before any AI project can start.' What is the BEST response to this objection?

- A) Agree and delay the AI initiative until the data warehouse is complete
- B) Override IT and launch the pilot without their involvement
- C) Start with use cases that work on unstructured text (emails, reports, documents) rather than requiring clean structured databases
- D) Hire external consultants to accelerate the data warehouse project to three months

Answer: C **Explanation:** C is correct because many high-value AI use cases — document summarisation,

email drafting, report analysis — operate on unstructured text and do not require a clean data warehouse. This sidesteps the objection without dismissing it. A is wrong because it allows a legitimate concern to become an indefinite blocker. B is wrong because alienating IT creates political opposition that will undermine the initiative. D is wrong because it accepts the false premise that a data warehouse must precede all AI use.

Q13. In a 90-day plan, when should the AI tool selection decision be made?

- A) After all pilots have concluded and results are compared
- B) Within the first two weeks, to prevent analysis paralysis from stalling the initiative
- C) Only after the legal department completes a full vendor risk assessment
- D) The tool decision should be delegated entirely to the IT department

Answer: B **Explanation:** B is correct because vendor evaluation paralysis is one of the top three reasons 90-day plans stall; making a decision quickly and adjusting later preserves momentum. A is wrong because waiting until after pilots conclude means no pilots were launched. C is wrong because a full vendor risk assessment takes weeks to months; a reasonable due-diligence check is sufficient for a pilot. D is wrong because tool selection should involve business stakeholders who understand the use case, not just IT.

Q14. Scenario: *A healthcare network's AI pilot for clinical note summarisation was led by a junior project manager with no budget authority. When the pilot needed to upgrade from a free AI tool to an enterprise version costing \$15,000, the request sat in budget approval for three weeks. By the time it was approved, the pilot team had lost momentum and two key participants had moved on to other projects. What was the PRIMARY root cause of this failure?*

- A) The AI tool was too expensive for a pilot project
- B) The pilot owner lacked the authority to clear budget and political obstacles, indicating insufficient executive sponsorship
- C) The pilot should have continued using the free AI tool to avoid budget complications
- D) Three weeks is a reasonable timeline for budget approval and the team should have been more patient

Answer: B **Explanation:** B is correct because the core issue was that the pilot owner lacked authority and there was no executive sponsor to expedite the budget request. A is wrong because \$15,000 is modest for enterprise AI. C is wrong because free tools lack data-protection guarantees needed for clinical data. D is wrong because three weeks of stall time on a 90-day plan represents over 20% of the entire timeline.

Q15. What is the relationship between a 90-day AI plan and a 12-month AI roadmap?

- A) They are the same document viewed at different levels of detail
- B) The 90-day plan produces the evidence and data needed to build a credible 12-month roadmap
- C) The 12-month roadmap must be completed before the 90-day plan can begin
- D) The 90-day plan replaces the need for a 12-month roadmap entirely

Answer: B **Explanation:** B is correct because the 90-day plan generates real pilot data — what worked,

what ROI was achieved, what obstacles arose — which grounds the 12-month roadmap in evidence rather than speculation. A is wrong because they are distinct documents with different purposes. C is wrong because waiting for a 12-month roadmap would defeat the urgency of the 90-day approach. D is wrong because organisations still need long-term strategic direction; the 90-day plan is the on-ramp, not the destination.

Q16. *Scenario:* An energy company's AI pilot programme has three pilots running: automated invoice processing (achieving 45% time reduction), predictive maintenance alerts (achieving 22% reduction in unplanned downtime), and AI-assisted contract review (achieving mixed results with 60% accuracy). The team is at week 8 and must decide what to scale. Applying the framework's scaling criteria, which pilot(s) should be prioritised for scaling?

- A) All three, to demonstrate broad AI capability across the organisation
- B) Only automated invoice processing, because it exceeds the 30% efficiency threshold
- C) Automated invoice processing and predictive maintenance, because both show measurable improvement; contract review needs refinement before scaling
- D) None — all pilots should run for another quarter before making scaling decisions

Answer: C **Explanation:** C is correct because invoice processing (45% improvement) clearly exceeds the 30% threshold, and predictive maintenance (22% reduction in unplanned downtime) represents significant value in an industry where unplanned downtime costs millions. Contract review at 60% accuracy needs further tuning. A is wrong because scaling a 60%-accuracy tool would damage credibility. B is wrong because it overlooks the high value of predictive maintenance in the energy sector. D is wrong because delaying decisions for another quarter violates the 90-day urgency principle.

Q17. When framing a budget request for AI pilots to senior leadership, what positioning is MOST effective?

- A) Emphasise the low cost relative to annual IT spend
- B) Frame the investment as risk mitigation against competitive displacement rather than an expense
- C) Request the budget as part of the R&D innovation fund to avoid scrutiny
- D) Avoid specifying a dollar amount and instead ask for 'resources as needed'

Answer: B **Explanation:** B is correct because executives respond to strategic risk framing ('Can we afford to be the last company without AI capabilities?') more than to cost-minimisation arguments. A is less compelling because it positions AI as a minor experiment rather than a strategic imperative. C is wrong because hiding AI investment under R&D reduces executive visibility and sponsorship. D is wrong because vague budget requests signal lack of planning and are more likely to be rejected.

Q18. *Scenario:* A technology company's new Chief Digital Officer creates a 90-day AI plan but keeps it confidential, sharing it only with her direct reports. She reasons that if the pilots fail, no one else needs to know. After 60 days, two of three pilots show strong results, but she struggles to get cooperation from other departments for the scaling phase. What was her CRITICAL strategic mistake?

- A) She should have started with more than three pilots to increase the chance of success
- B) She failed to build broad organisational awareness and buy-in, making cross-department scaling politically difficult
- C) She should have waited for all three pilots to succeed before approaching other departments
- D) She should have outsourced the scaling phase to an external consulting firm

Answer: B **Explanation:** B is correct because keeping the initiative confidential prevented the communication and relationship-building needed for cross-department cooperation during scaling. Public commitment and visibility are essential. A is wrong because more pilots would have further strained her limited team. C is wrong because waiting for perfection is unnecessary; two strong results is sufficient. D is wrong because external consultants cannot substitute for internal political capital.

Q19. Which of the following is the BEST example of a specific, completable 'tomorrow action' in the personal commitments section?

- A) 'Start thinking about AI use cases for my department'
- B) 'Send an email to the VP of Customer Service proposing a 15-minute meeting about the AI pilot on Thursday'
- C) 'Research AI tools online when I have time'
- D) 'Discuss AI with the team at the next staff meeting'

Answer: B **Explanation:** B is correct because it is specific (email, named recipient), completable in one action, and time-bound (Thursday meeting). A is wrong because 'start thinking' is not an action. C is wrong because 'when I have time' ensures it never happens. D is wrong because 'next staff meeting' may be weeks away; tomorrow actions must be immediate.

Q20. *Scenario:* A consulting firm's managing partner completed the 90-day workshop template but has not looked at it since. Two weeks later, no actions have been taken. The firm's competitors have started offering AI-powered deliverables to clients. What was the MOST likely reason the plan was not executed?

- A) The plan was too ambitious and should have been scaled back to a 30-day plan
- B) The managing partner did not share the plan or establish an accountability mechanism, allowing it to be crowded out by daily operations
- C) The competitive threat was not urgent enough to justify action
- D) The 90-day framework does not work for professional services firms

Answer: B **Explanation:** B is correct because the absence of accountability (sharing with others, calendaring follow-ups) is the primary predictor of plan abandonment. A is wrong because the plan's length was appropriate; the problem was execution, not scope. C is wrong because competitors offering AI-powered deliverables is an immediate threat in consulting. D is wrong because the framework is industry-agnostic.