

Human-in-the-Loop Protocol

Generative AI for Business Leaders & C-Suite | AI Governance & Oversight

1. Why Humans Still Matter in the Age of AI

Artificial intelligence excels at processing volume, spotting statistical patterns, and executing repetitive tasks at speeds no human team can match. What it cannot do is exercise moral judgment, read the political subtext of a boardroom decision, or sense that an employee's resignation letter carries deeper organisational warning signs. Human-in-the-loop (HITL) governance exists because AI's strengths and weaknesses are the mirror image of human strengths and weaknesses. Combining both produces decisions that are faster, more scalable, and safer than either could achieve alone.

Think of HITL like the relationship between an autopilot system and a commercial airline pilot. The autopilot handles cruising altitude, minor course corrections, and routine adjustments far more consistently than any human. But takeoff, landing, severe turbulence, and engine failures still require a pilot in command. Nobody boards a plane expecting zero human oversight, and no organisation should deploy AI in consequential decision-making without the same principle.

Key Point: Human-in-the-loop is not a brake on AI's potential. It is a design principle that makes AI safe enough to deploy at scale. Organisations that skip HITL governance do not move faster — they move toward their first catastrophic error faster.

2. The Three-Tier Review Framework

The most practical way to implement HITL is to match the level of human oversight to the stakes of the decision. A three-tier model prevents two equally dangerous mistakes: under-reviewing high-stakes AI outputs (risk of harm) and over-reviewing low-stakes outputs (risk of bottlenecks that make people abandon the process entirely).

Tier	Decision Stakes	Human Oversight Required	Examples
Tier 1 High-Stakes	Decisions affecting people's lives, livelihoods, legal standing, or major financial commitments	Expert human review and approval before every action. No exceptions.	Hiring/firing decisions, credit approvals over a threshold, regulatory filings, medical/safety determinations, M&A strategy
Tier 2 Medium-Stakes	Decisions with moderate business or reputational impact	Spot-check sampling (10-20% of outputs) plus periodic quality reviews	Customer complaint responses, marketing content for publication, internal policy changes, moderate-value transactions
Tier 3 Low-Stakes	Routine operational outputs with limited external impact	Monitoring dashboards and quarterly audits, but no pre-approval gate	Internal meeting summaries, routine email drafts, data analysis for internal use, administrative automation

The tier assignment is not permanent. A task that normally sits at Tier 3 — say, generating internal meeting notes — escalates to Tier 1 if the meeting involves sensitive merger discussions or personnel decisions. Context determines the tier, not the task category alone.

3. What Effective Human Review Actually Looks Like

Putting a human in the loop is meaningless if that human is rubber-stamping outputs without genuine scrutiny. Effective review requires training, tools, authority, and accountability.

Training reviewers on AI failure modes

Most AI errors follow predictable patterns. Reviewers who know these patterns catch problems far more reliably than reviewers told simply to 'check the output.' The most common failure modes include:

- **Hallucination:** The AI states a fact, cites a source, or quotes a statistic that does not exist. Large language models generate plausible-sounding text, not verified truth. Financial figures, legal precedents, and scientific citations are especially vulnerable.
- **Bias amplification:** AI trained on historical data reproduces and sometimes amplifies the biases embedded in that data. A hiring model trained on ten years of predominantly male engineering hires will systematically downrank female candidates unless explicitly corrected.
- **Logical inconsistency:** The AI contradicts itself within the same output — for example, recommending budget cuts in paragraph two after recommending increased investment in paragraph one.
- **Tone-deafness:** AI lacks emotional intelligence. It may generate a cheerful, enthusiastic response to a customer reporting a data breach, or use casual language in a regulatory filing.

Structured evaluation criteria

Do not ask reviewers to 'use their judgment' without structure. Provide a review checklist or rubric that operationalises quality. A simple four-question framework works well for most business contexts:

Review Question	What It Catches
Is every factual claim verifiable against source data?	Hallucinations, fabricated statistics, invented citations
Would this output cause harm or embarrassment if it reached the wrong audience?	Tone errors, confidentiality leaks, reputational risks
Does the output treat all individuals and groups equitably?	Bias, discriminatory language, exclusionary framing
Is the reasoning internally consistent from start to finish?	Logical contradictions, conflicting recommendations

Reviewer authority and tracking

Reviewers must have genuine authority to reject, modify, or return AI outputs. If the organisational culture penalises reviewers for slowing down production, the review process becomes performative. Track rejection rates, modification rates, and the types of issues caught. If rejection rates are near zero,

the review is almost certainly too lenient. If they are extremely high, the AI tool may need retraining or replacement.

4. Escalation Triggers — When to Override the Tier System

Even well-designed tier systems need an override mechanism. Certain situations demand full Tier 1 human review regardless of the task's normal classification:

- **Novel situations:** AI performs worst on tasks it has never encountered. If the business context is genuinely unprecedented — a new market, a regulatory change, an emerging technology — human judgment is essential.
- **Emotionally sensitive contexts:** Layoff communications, bereavement responses, crisis PR. AI can draft, but a human must own the emotional calibration.
- **Ethical grey areas:** Decisions where reasonable people disagree on the right course of action. AI cannot weigh competing values; it optimises for whatever metric it was trained on.
- **High-visibility situations:** Anything that could appear in the press, in a regulatory inquiry, or in front of a board of directors. The reputational cost of an AI error here dwarfs any efficiency gain.
- **Adversarial contexts:** Situations where someone may deliberately try to manipulate or exploit the AI — prompt injection attacks, social engineering via AI chatbots, or competitive intelligence probing.

Real-World Incident — Air Canada Chatbot (2024): Air Canada's AI chatbot told a grieving customer he could book a full-fare flight and later apply for a bereavement discount. This was false — the airline's actual policy required the discount to be applied at the time of booking. When the customer demanded the promised discount, Air Canada argued the chatbot was a 'separate legal entity' responsible for its own statements. A Canadian tribunal ruled against the airline, holding that Air Canada was liable for its chatbot's representations regardless of whether a human approved them. The case established that deploying an AI in a customer-facing role without adequate human oversight does not shield the organisation from liability for the AI's errors.

5. The Human-AI Partnership Model

The goal of HITL governance is not 'human versus AI' — it is designing workflows where each handles what it does best. This is the partnership principle.

AI Handles	Human Handles	Combined Outcome
Speed and scale — processing thousands of documents, transactions, or data points simultaneously	Context and judgment — understanding what the numbers mean for this specific business situation	Faster decisions that account for nuance
Pattern recognition — identifying statistical trends, anomalies, and correlations across large datasets	Ethical reasoning — deciding what should be done, not just what can be done	Data-driven decisions with moral guardrails
First-draft generation — producing initial analysis, reports, or communications	Refinement and sign-off — ensuring accuracy, tone, and strategic alignment	Higher-quality output in less total time

Continuous monitoring — watching dashboards, flagging exceptions, tracking KPIs 24/7	Strategic interpretation — deciding when a flagged anomaly requires action versus investigation versus ignoring	Reliable surveillance with intelligent response
--	---	---

Research consistently shows that humans working with AI outperform either humans or AI working alone. A 2023 study by the Boston Consulting Group found that consultants using GPT-4 completed tasks 25% faster and produced 40% higher-quality outputs than consultants working without AI — but only when they understood the AI's limitations and maintained critical oversight. Consultants who blindly accepted AI outputs actually produced worse results than those who worked without AI at all.

6. Common HITL Implementation Mistakes

Organisations frequently undermine their own HITL frameworks through predictable errors. Awareness of these patterns helps leaders design systems that actually work.

- **Automation bias:** Humans tend to trust computer-generated outputs more than they should, especially when the output looks polished and confident. Reviewers must be trained to treat AI outputs as hypotheses to verify, not conclusions to accept.
- **Review fatigue:** If reviewers are asked to check hundreds of AI outputs daily, quality deteriorates rapidly. Rotate reviewers, limit daily review volume, and automate the easy checks so humans focus on genuinely ambiguous cases.
- **Missing feedback loops:** When reviewers catch an error, the correction often stays with the individual case rather than feeding back into the AI system. Build processes that route reviewer corrections into model retraining or prompt improvement.
- **Insufficient reviewer authority:** If rejecting an AI output creates friction — extra paperwork, manager pushback, missed deadlines — reviewers will stop rejecting. Make rejection easy and celebrated, not penalised.

7. Building a HITL Process — Step by Step

For leaders implementing HITL governance from scratch, follow this sequence:

- **Step 1 — Inventory all AI-assisted decisions.** List every process where AI influences or generates outputs. Many organisations discover AI tools being used informally that were never formally approved.
- **Step 2 — Classify each decision by tier.** Use the three-tier framework. Involve legal, compliance, and business owners in the classification. Err on the side of higher tiers initially; you can relax later with evidence.
- **Step 3 — Design the review workflow for each tier.** Define who reviews, using what criteria, within what timeframe, with what authority to reject.
- **Step 4 — Train your reviewers.** Cover AI failure modes, your evaluation criteria, and the escalation triggers. Training is not optional.
- **Step 5 — Instrument and measure.** Track rejection rates, error types caught, review turnaround times, and user compliance. Data tells you whether the system is working.

- **Step 6 — Review and adjust quarterly.** AI capabilities change, your use cases evolve, and your reviewers learn. The HITL framework should evolve with them.

8. Metrics That Matter for HITL Governance

Metric	What It Tells You	Warning Sign
Review rejection rate	How often reviewers override AI outputs	Near 0% (rubber-stamping) or above 50% (AI tool not fit for purpose)
Time-to-review	How long human review adds to the workflow	Consistently exceeding SLA, causing teams to skip review
Error escape rate	Proportion of AI errors that reach end users despite review	Any upward trend; indicates reviewer fatigue or inadequate training
Escalation frequency	How often tier overrides are triggered	Very low (people not escalating when they should) or very high (tier assignments too low)
Reviewer satisfaction	Whether reviewers feel supported and empowered	Low scores correlate with declining review quality

9. Key Terms Glossary

Term	Definition
Human-in-the-Loop (HITL)	A governance model in which human judgment is embedded at defined points in an AI-assisted workflow, ensuring that consequential decisions receive human oversight before execution.
Automation Bias	The tendency for humans to over-rely on automated or computer-generated outputs, accepting them uncritically even when they contain errors.
Tier Classification	A system for categorising AI-assisted decisions by their potential impact (high, medium, low) to determine the appropriate level of human oversight.
Escalation Trigger	A predefined condition or context that requires a decision to receive higher-tier human review than its default classification would normally require.
Review Rubric	A structured set of evaluation criteria that reviewers apply consistently when assessing AI-generated outputs, replacing subjective judgment with systematic checks.
Error Escape Rate	The proportion of AI-generated errors that pass through human review undetected and reach the end user or downstream process.
Feedback Loop	A process by which reviewer corrections and error findings are routed back to improve the AI model, prompts, or workflow configuration.
Adversarial Context	A situation in which an external party deliberately attempts to manipulate, exploit, or extract information from an AI system

	through crafted inputs.
--	-------------------------

10. Quick Revision Checklist

- AI augments human decision-making — it does not replace it for consequential decisions.
- The three-tier framework matches oversight intensity to decision stakes: Tier 1 (always review), Tier 2 (spot-check 10-20%), Tier 3 (monitor and audit).
- Tier assignments are context-dependent — the same task can escalate based on sensitivity, novelty, or visibility.
- Effective review requires trained reviewers, structured criteria, access to source data, and genuine authority to reject.
- Track rejection rates: near-zero signals rubber-stamping; very high signals a broken AI tool.
- Five escalation triggers override normal tier assignments: novel situations, emotional contexts, ethical grey areas, high visibility, and adversarial contexts.
- The human-AI partnership principle: AI handles speed, scale, and pattern recognition; humans handle judgment, ethics, and accountability.
- Automation bias is the single biggest threat to HITL effectiveness — train against it explicitly.
- Build feedback loops so reviewer corrections improve the AI system, not just the individual case.
- Review and adjust the HITL framework quarterly as AI capabilities and use cases evolve.

20 Exam-Based MCQs — Human-in-the-Loop Protocol

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. Which tier of the three-tier review framework requires expert human review before EVERY action?

- A) Tier 3 — Low-stakes decisions
- B) Tier 1 — High-stakes decisions
- C) Tier 2 — Medium-stakes decisions
- D) All tiers require review before every action

Answer: B **Explanation:** B is correct because Tier 1 covers high-stakes decisions affecting employment, major finances, legal matters, safety, or strategy, and requires no exceptions to the review rule. A is wrong because Tier 3 uses monitoring and audits, not pre-approval. C is wrong because Tier 2 uses spot-check sampling. D is wrong because requiring review on every output across all tiers would create unsustainable bottlenecks.

Q2. *Scenario:* A marketing manager uses AI to generate a social media post for a routine product update. The company's HITL framework classifies routine marketing content as Tier 2 (medium-stakes). However, the post references a competitor by name and makes performance comparison claims. What should the marketing manager do?

- A) Post it — Tier 2 only requires spot-check sampling, and this post may not be sampled
- B) Escalate to Tier 1 review — competitor comparisons create legal and reputational risk that overrides the default tier
- C) Delete the competitor reference and post without review
- D) Ask the AI to verify its own claims about the competitor

Answer: B **Explanation:** B is correct because competitor comparisons carry legal risk (potential defamation, false advertising) and reputational risk, triggering the 'high-visibility' and 'ethical grey area' escalation criteria. A is wrong because following the default tier ignores the escalation triggers clearly present. C is wrong because removing one element does not address whether other claims need review. D is wrong because AI cannot verify its own factual claims — this is precisely the kind of task requiring human expert review.

Q3. What is automation bias?

- A) The tendency for AI systems to prefer automated processes over manual ones
- B) The tendency for humans to over-rely on computer-generated outputs, accepting them uncritically
- C) A programming error that causes AI to automate the wrong tasks
- D) The preference organisations have for fully automated solutions

Answer: B **Explanation:** B is correct because automation bias is a well-documented cognitive bias in which humans treat computer outputs as more trustworthy than they are, reducing critical scrutiny. A is wrong because automation bias describes human behaviour, not AI behaviour. C is wrong because it describes a technical bug, not a cognitive bias. D is wrong because it describes an organisational preference, not the specific psychological phenomenon.

Q4. Scenario: A financial services firm has implemented HITL review for AI-generated investment research. After six months, the review rejection rate is 0.3% — reviewers override AI outputs in fewer than 1 in 300 cases. Management is pleased with the 'high accuracy' of the AI. What concern should a governance officer raise about this rejection rate?

- A) The rate is appropriate — it demonstrates the AI is performing well
- B) The rate is too high — reviewers are being too cautious
- C) The near-zero rate likely indicates rubber-stamping rather than genuine review
- D) The firm should eliminate human review since the AI is clearly accurate enough

Answer: C **Explanation:** C is correct because a rejection rate near zero is a classic warning sign of automation bias — reviewers are likely approving outputs without genuine scrutiny. No AI system achieves 99.7% accuracy on complex financial research. A is wrong because interpreting near-zero rejection as 'accuracy' ignores the strong probability of inadequate review. B is wrong because 0.3% is extremely low, not too high. D is wrong because removing review based on potentially misleading metrics would eliminate the safety net entirely.

Q5. In the three-tier framework, what percentage of outputs should typically be spot-checked at Tier 2?

- A) 1-5%
- B) 10-20%
- C) 50-75%
- D) 100%

Answer: B **Explanation:** B is correct because the framework recommends 10-20% sampling for medium-stakes decisions, balancing thoroughness with operational efficiency. A is wrong because 1-5% is too low to catch systematic issues. C is wrong because 50-75% approaches Tier 1 levels of oversight, defeating the purpose of tiering. D is wrong because 100% review is Tier 1, not Tier 2.

Q6. *Scenario:* A healthcare organisation uses AI to generate initial patient appointment summaries for doctors. Normally classified as Tier 3 (low-stakes internal documentation), the system is now being asked to summarise notes from psychiatric consultations involving suicidal ideation. How should the HITL framework respond?

- A) Maintain Tier 3 — the task is still appointment summaries
- B) Disable AI for psychiatric consultations entirely
- C) Escalate to Tier 1 — the content involves safety-critical and emotionally sensitive information
- D) Increase the Tier 3 audit frequency from quarterly to monthly

Answer: C **Explanation:** C is correct because psychiatric notes involving suicidal ideation trigger multiple escalation criteria: safety-critical content, emotionally sensitive context, and potential for serious harm if the AI misrepresents or omits critical details. A is wrong because the task category does not override the sensitivity of the content. B might be considered but is overly restrictive — AI can assist if properly reviewed. D is wrong because increasing audit frequency does not address the immediate risk of each individual summary.

Q7. Which of the following is NOT one of the five escalation triggers that override normal tier classification?

- A) Novel situations where AI has not been tested
- B) Budget constraints that limit available reviewers
- C) High-visibility situations that could attract media attention
- D) Adversarial contexts where someone may try to exploit the AI

Answer: B **Explanation:** B is correct because budget constraints are a resource limitation, not a decision characteristic that increases risk. The five escalation triggers are: novel situations, emotionally sensitive contexts, ethical grey areas, high-visibility situations, and adversarial contexts. A, C, and D are all legitimate escalation triggers.

Q8. *Scenario:* A regional bank has deployed AI to pre-screen mortgage applications. The AI recommends approval or denial, and a loan officer reviews each recommendation. After three months, the loan officer

reports that she reverses the AI's recommendation in about 15% of cases. Her manager tells her she is 'slowing down the process' and should 'trust the system more.' What is the MOST significant governance risk in this situation?

- A) The AI's 85% agreement rate suggests it needs retraining
- B) The manager is undermining reviewer authority, which will lead to rubber-stamping
- C) The loan officer is being too cautious and should reduce her rejection rate
- D) The bank should replace human review with automated verification

Answer: B **Explanation:** B is correct because pressuring reviewers to accept more AI outputs directly undermines the HITL principle. If the loan officer feels penalised for exercising judgment, she will stop pushing back — turning review into rubber-stamping. A is wrong because 85% agreement is not inherently a problem without knowing what errors exist. C is wrong because 15% reversal is a reasonable rate that suggests genuine review. D is wrong because mortgage decisions are high-stakes (Tier 1) and require human judgment on legal, financial, and fairness grounds.

Q9. In the human-AI partnership model, which function is best handled by humans rather than AI?

- A) Processing thousands of transactions simultaneously
- B) Monitoring dashboards and flagging exceptions 24/7
- C) Deciding what should be done versus what can be done
- D) Identifying statistical patterns across large datasets

Answer: C **Explanation:** C is correct because ethical reasoning — distinguishing what should be done from what can be done — requires human judgment, values, and contextual understanding that AI lacks. A, B, and D are all tasks where AI excels due to speed, scale, and pattern recognition capabilities.

Q10. *Scenario:* A consulting firm has trained its AI review team to check for hallucinations, bias, and tone issues. After implementing structured review checklists, the team catches errors consistently. However, when reviewers flag problems, the corrections are applied to the individual document but never fed back to improve the AI system's prompts or configuration. What governance gap does this represent?

- A) Insufficient reviewer training
- B) Missing feedback loops between review findings and AI improvement
- C) Inadequate tier classification
- D) Automation bias among reviewers

Answer: B **Explanation:** B is correct because the firm has good detection but no mechanism to route corrections back into the AI system. Without feedback loops, the same errors recur, and reviewers must catch them repeatedly. A is wrong because the reviewers are catching errors effectively. C is wrong because the issue is post-review process, not classification. D is wrong because the reviewers are actively catching errors, not rubber-stamping.

Q11. Why is the airline autopilot analogy useful for understanding human-in-the-loop governance?

- A) It shows that AI should eventually replace all human involvement

- B) It illustrates that routine tasks can be automated while critical moments require human control
- C) It proves that AI is always more reliable than humans
- D) It demonstrates that human oversight is only needed during emergencies

Answer: B **Explanation:** B is correct because the autopilot analogy shows that automation handles routine operations well (cruising altitude, minor corrections) while humans remain essential for high-stakes situations (takeoff, landing, emergencies). A is wrong because the analogy explicitly retains the pilot. C is wrong because the analogy acknowledges both strengths. D is wrong because human pilots are involved in takeoff and landing (planned events), not just emergencies.

Q12. Scenario: An e-commerce company's AI customer service chatbot handles 80% of customer inquiries automatically. The remaining 20% are escalated to human agents. Recently, a customer asked the chatbot about the company's return policy for a product that was recalled due to a safety defect. The chatbot provided the standard return policy without mentioning the recall or safety issue. Which HITL failure does this scenario illustrate?

- A) Automation bias among customers
- B) Inadequate tier classification — safety-related inquiries should trigger automatic escalation to human agents
- C) Review fatigue among the human agents
- D) The chatbot needs more training data about recalls

Answer: B **Explanation:** B is correct because safety-related inquiries are a clear escalation trigger (safety context, potential liability, reputational risk) that should automatically route to a human agent, not be handled by the chatbot at all. A is wrong because the issue is the system design, not customer behaviour. C is wrong because the human agents never saw this inquiry. D is wrong because the fundamental issue is governance design, not model training — the chatbot should not be handling safety-critical inquiries regardless of its training.

Q13. What is the recommended approach when an AI-assisted decision falls into an ethical grey area?

- A) Let the AI decide since it is objective and unbiased
- B) Apply the default tier classification since the framework covers all situations
- C) Escalate to Tier 1 (full human review) regardless of the decision's normal tier
- D) Defer the decision until AI technology improves

Answer: C **Explanation:** C is correct because ethical grey areas are one of the five escalation triggers that override normal tier classification. When reasonable people disagree on the right course, human judgment and values must prevail. A is wrong because AI is not objective — it optimises for its training objective, which may not align with ethical considerations. B is wrong because escalation triggers exist precisely to override defaults. D is wrong because deferring avoids the responsibility to make a governed decision now.

Q14. Scenario: A pharmaceutical company uses AI to draft adverse event reports for regulatory submission. A quality reviewer notices that the AI occasionally misclassifies the severity of patient side

effects — calling 'serious' events 'non-serious' in about 2% of cases. The reviewer raises the issue, but management decides the 98% accuracy rate is acceptable. Why is management's position problematic?

- A) 98% accuracy is below the industry standard of 99.5%
- B) Misclassifying adverse event severity in regulatory filings creates legal liability and patient safety risks, regardless of the error rate
- C) The reviewer should have caught all errors before raising the issue
- D) The AI should be replaced with a different vendor's product

Answer: B **Explanation:** B is correct because adverse event reporting is a Tier 1 (high-stakes) function where any misclassification can result in regulatory action, patient harm, and legal liability. A 2% error rate on severity classification is not a statistical accuracy question — it is a patient safety and compliance question. A is wrong because the issue is not about hitting a numerical threshold but about the nature of the errors. C is wrong because the reviewer did the right thing by escalating a systemic issue. D is wrong because replacing the vendor does not address the governance gap.

Q15. What is review fatigue and how should organisations mitigate it?

- A) Fatigue caused by too few reviews; mitigate by increasing review volume per person
- B) Fatigue from reviewing too many AI outputs, leading to declining quality; mitigate by rotating reviewers and limiting daily volume
- C) Fatigue caused by reviewing outputs that are always correct; mitigate by adding artificial errors
- D) Fatigue from using multiple review tools; mitigate by standardising on one platform

Answer: B **Explanation:** B is correct because review fatigue occurs when reviewers check too many outputs and their attention and quality deteriorate. Mitigation includes rotating reviewers, capping daily review counts, and automating easy checks so humans focus on genuinely ambiguous cases. A is wrong because it inverts the problem. C is wrong because it describes a different issue and an impractical solution. D is wrong because tool consolidation addresses a workflow problem, not a cognitive fatigue problem.

Q16. *Scenario:* A law firm uses AI to draft initial contract reviews. Junior associates review the AI's output before sending to partners. The firm tracks that associates modify AI drafts in 30% of cases. A new managing partner proposes reducing the associate review step to 'save costs and improve turnaround time.' What is the strongest argument against removing associate review?

- A) Associates need the training experience of reviewing contracts
- B) A 30% modification rate indicates the AI produces errors in nearly one-third of outputs — removing review would send flawed analysis to partners and clients
- C) Removing review would violate employment law
- D) The AI vendor's terms of service require human review

Answer: B **Explanation:** B is correct because a 30% modification rate demonstrates that associates are providing substantive corrections in a significant portion of outputs. Removing this layer would expose the firm to errors in client-facing legal work — a Tier 1 function with serious liability implications. A has some merit but is not the strongest governance argument. C is incorrect as a factual claim. D may or may not be true but is not the core governance concern.

Q17. According to the BCG study referenced, what happened when consultants blindly accepted AI outputs without maintaining critical oversight?

- A) They produced results 25% faster than those working without AI
- B) They produced worse results than consultants who worked without AI at all
- C) They produced equivalent results to those using AI with oversight
- D) They produced slightly better results than working without AI

Answer: B **Explanation:** B is correct because the 2023 BCG study found that consultants who uncritically accepted AI outputs actually performed worse than those working without AI — demonstrating that blind trust in AI is worse than no AI at all. A is wrong because that speed improvement applied to consultants who used AI with proper oversight. C and D are wrong because they understate the negative impact of uncritical AI acceptance.

Q18. *Scenario:* A retail company's AI generates weekly sales forecasts. The data analytics team reviews forecasts before they go to regional managers. The team lead notices that one analyst consistently approves every AI forecast without modification, even during weeks with unusual market conditions (holiday periods, supply disruptions). Other analysts modify forecasts 10-15% of the time. What should the team lead do FIRST?

- A) Remove the analyst from the review rotation immediately
- B) Investigate whether the analyst is exhibiting automation bias and provide additional training on AI failure modes and review expectations
- C) Accept the analyst's reviews since the AI is probably accurate
- D) Implement automated quality checks to replace human review

Answer: B **Explanation:** B is correct because the zero-modification pattern during unusual market conditions is a strong indicator of automation bias. The first step is investigation and training, not punishment. A is wrong because removing the analyst without understanding the cause does not fix the underlying issue. C is wrong because accepting uncritical reviews, especially during anomalous periods, defeats the purpose of HITL. D is wrong because automated checks complement but do not replace human review for contextual judgment.

Q19. What is the FIRST step in building a HITL process from scratch?

- A) Train all employees on AI failure modes
- B) Inventory all AI-assisted decisions across the organisation
- C) Design review workflows for each decision type
- D) Implement monitoring dashboards

Answer: B **Explanation:** B is correct because you cannot design oversight for processes you have not identified. The inventory step often reveals informal AI usage that was never formally approved — a critical discovery for governance. A is important but comes after you know what needs reviewing. C requires the inventory and tier classification to be complete first. D is an instrumentation step that comes after the

framework is designed.

Q20. Scenario: A government agency is deploying AI to help process benefits applications. The agency designs a HITL framework where case workers review all AI recommendations before decisions are issued. After the first month, case workers report that the AI's recommendation screen does not show the underlying data or reasoning — only the recommended decision (approve/deny) and a confidence score. What is the MOST critical flaw in this HITL implementation?

- A) The confidence score should be removed to prevent bias
- B) Case workers cannot verify AI claims against source data, making meaningful review impossible
- C) The AI should make final decisions without case worker review to improve speed
- D) The agency should increase the number of case workers to handle the volume

Answer: B Explanation: B is correct because effective review requires access to source material so reviewers can verify AI claims against original data. Showing only a decision and confidence score forces reviewers into rubber-stamping because they have no basis for independent judgment. A is wrong because removing the confidence score does not address the core problem. C is wrong because government benefits decisions are Tier 1 and require human oversight. D is wrong because more reviewers cannot compensate for a system that does not provide the information needed for meaningful review.