Interactive Exercise: Your AI Opportunity Assessment

Generative AI for Business Leaders & C-Suite | AI Fundamentals

1. Why Structured Opportunity Assessment Matters

Most AI initiatives fail not because the technology does not work, but because organisations pick the wrong problem to solve first. A 2024 Boston Consulting Group survey of fourteen hundred executives found that seventy-four percent of companies struggle to move AI projects beyond the pilot stage. The single biggest predictor of failure: choosing a use case that was either too ambitious for a first project or too trivial to demonstrate meaningful value.

A structured opportunity assessment prevents both mistakes. It forces you to evaluate potential AI applications against objective criteria — frequency, complexity, data availability, and business impact — rather than defaulting to whatever the loudest stakeholder suggests or whatever a vendor demo made look impressive. The exercise in this lecture gives you a repeatable framework you can use across every department in your organisation.

2. Step One: Identifying Your Time Drains

The exercise begins by asking you to list the three tasks that consume the most time for you or your team. This is not a random starting point — it is grounded in a principle from operations management: the highest-impact process improvements target the activities that consume the most resources relative to the value they produce.

The key discipline here is specificity. Writing 'reports' on your list tells you nothing actionable. Writing 'compiling the monthly regional sales performance report from five different data sources into a standardised PowerPoint template' tells you exactly what AI would need to do, what inputs it requires, and how you would measure success.

How to identify the right tasks

- Track your team's actual time allocation for one week — perception often differs from reality
- Focus on tasks that are repetitive across cycles (daily, weekly, monthly) rather than one-off projects
- Distinguish between the valuable core of a task and the overhead around it — AI often automates the overhead while humans retain the core
- Ask frontline staff, not just managers — they know where time disappears

> **Key Point:** The best AI pilot targets are hiding in plain sight: they are the tasks your team complains about doing, takes longer than it should, and follows a roughly similar pattern every time it is performed.

3. Step Two: Mapping Tasks to AI Capability Categories

The lecture identifies four AI capability categories. Understanding these categories is essential because they determine which AI tools and approaches apply to each task, and they help you communicate requirements to technical teams without needing to specify the underlying technology.

Category	What AI Does	Example Tools	Typical Business Tasks
Text Generation	Creates written content from instructions or context	ChatGPT, Claude, Gemini, Jasper	Emails, reports, proposals, documentation, marketing copy, meeting summaries
Data Analysis & Insights	Interprets data, finds patterns, produces visualisations	ChatGPT Advanced Data Analysis, Claude, Tableau AI	Financial analysis, sales trend identification, customer segmentation, KPI reporting
Research & Synthesis	Gathers information from multiple sources and summarises	Perplexity, Claude, ChatGPT with browsing	Competitive analysis, market research, document review, policy comparison
Creative Content	Generates visual, audio, or multimedia content	Midjourney, DALL-E, Canva AI, Runway	Marketing visuals, social media content, presentation graphics, product mockups

Many tasks span multiple categories. A quarterly board report, for example, involves data analysis (interpreting financial results), research (gathering competitive context), and text generation (writing the narrative). For your initial assessment, identify the primary category — the one that represents the bulk of the time spent. You can add AI capabilities for the secondary categories later.

4. Step Three: The Prioritisation Matrix

The lecture introduces a two-axis prioritisation framework based on frequency and expertise level. This is a simplified but effective model drawn from process improvement methodologies. The logic is straightforward: tasks that happen often and do not require deep specialist knowledge offer the fastest, lowest-risk returns from AI automation.

	Low-Medium Expertise Required	High Expertise Required
High Frequency (daily/weekly)	PRIORITY 1 — Start here. Quick wins, high volume, low risk. AI handles the repetitive work; minimal specialist oversight needed.	PRIORITY 2 — High value but needs expert-in-the-loop. AI drafts, expert reviews. Significant time savings with manageable risk.
Low Frequency (monthly/quarterly)	PRIORITY 3 — Moderate impact. Worth automating if the setup effort is low, but do not invest heavily in a task performed twelve times a year.	PRIORITY 4 — Park this. Complex, infrequent tasks are poor first AI projects. Revisit after your team has AI experience.

Why frequency matters more than you think

A task that takes thirty minutes and happens twice daily consumes two hundred and sixty hours per year per person. Reducing that by fifty percent with AI saves one hundred and thirty hours — equivalent to over three working weeks. The same thirty-minute task performed monthly saves only three hours per year. Frequency is the multiplier that turns modest per-instance savings into transformational annual impact.

Why expertise level matters

Expertise level is a proxy for risk. When AI handles a task that anyone on the team could verify (scheduling, data formatting, standard email responses), the cost of an AI error is low — someone catches it quickly. When AI handles a task that requires specialised knowledge to evaluate (legal analysis, medical interpretation, financial modelling), errors are harder to detect and potentially more costly. Starting with lower-expertise tasks lets your organisation build AI governance muscle before tackling high-stakes applications.

5. Building a Complete Opportunity Assessment

The lecture's exercise covers the essential first pass. For a thorough assessment that you can present to leadership or a steering committee, expand the framework to include additional evaluation criteria.

Criterion	Question to Answer	Why It Matters
Data availability	Does the task have accessible, digital input data?	AI needs data to work with. A task that relies on verbal conversations or physical documents requires a digitisation step before AI can help.
Process consistency	Is the task performed roughly the same way each time?	High consistency means clearer patterns for AI to learn. Highly variable processes are harder to automate.
Error tolerance	What happens if AI gets it wrong?	Low error tolerance (medical, legal, financial) requires human-in-the-loop. High error tolerance (internal drafts, brainstorming) allows more autonomy.
Measurability	Can you quantify the current cost and the improvement?	You need a baseline to prove ROI. If you cannot measure the current state, you cannot demonstrate the value AI adds.
Stakeholder readiness	Will the people involved accept AI assistance?	Technical feasibility without stakeholder buy-in leads to shelfware. Include change management in your assessment.

6. From Assessment to Action: The Seven-Day Challenge

The lecture closes with a commitment to test AI on your priority task within seven days. This is not arbitrary motivation — it is based on a well-documented principle in change management: the faster people experience a new technology personally, the faster they move from scepticism to adoption.

How to run your seven-day test effectively

- **Day 1-2: Set up your tool.** Choose ChatGPT, Claude, or Gemini. Create an account if you do not have one. Spend thirty minutes exploring the interface.

- **Day 3-4: Run three real examples.** Take three actual instances of your priority task and complete them with AI assistance. Time both the AI-assisted and traditional approaches.
- **Day 5-6: Evaluate quality.** Compare AI outputs against what you or your team would have produced manually. Note where AI was faster, where it was as good, and where it fell short.
- **Day 7: Document and decide.** Write a one-page summary: time saved, quality comparison, surprises, and a recommendation on whether to scale this to the team.

Real-World Incident — Deloitte's AI Experimentation Programme (2023): Deloitte launched an internal programme encouraging consultants to use AI tools on real client deliverables — under supervision — for two weeks. The result: eighty-two percent of participants reported that AI saved them meaningful time on at least one task type, and the most common reaction was 'I didn't realise it could do that.' Personal experimentation converted more sceptics than any executive presentation or vendor demo could.

7. Common Pitfalls in AI Opportunity Assessment

Starting too big

Leaders often want their first AI project to be transformational — a complete reimagining of customer service or a fully automated supply chain. These projects take months to deliver, involve dozens of stakeholders, and create political pressure to declare success before real results emerge. Start with a task one person can test in one week.

Choosing a task nobody cares about

The opposite mistake is picking something so minor that even a successful outcome generates no enthusiasm. Automating the formatting of an internal memo template saves five minutes a week and impresses nobody. Find the sweet spot: a task that is painful enough that people will notice the improvement, but simple enough that AI can deliver results quickly.

Confusing tool capability with business readiness

Just because AI can do something does not mean your organisation is ready for it. If your customer data is scattered across seven systems with no integration, AI-powered customer personalisation is a data engineering project first and an AI project second. Assess the prerequisites, not just the endpoint.

Skipping measurement

If you do not establish a baseline before introducing AI, you cannot prove its value. Measure the current time, cost, and quality of the task before you start. Then measure the same metrics with AI. Without this comparison, your AI project becomes an anecdote instead of a business case.

8. Scaling Beyond the First Use Case

Once your seven-day pilot succeeds, the natural question is: what next? The scaling path follows a predictable pattern:

- **Same task, more people:** If AI worked for you, roll it out to the rest of the team doing the same task. This is the lowest-risk expansion.

- **Adjacent tasks, same function:** If AI helped with sales email drafting, try it for proposal writing, call preparation, and meeting summaries within the sales function.
- **Same capability, different function:** If text generation worked in sales, try it in HR (job descriptions), legal (contract summaries), and finance (management commentary).
- **Higher-complexity tasks:** Once the team has AI fluency and governance processes, move to Priority 2 and Priority 3 tasks from your matrix.

9. Building the Business Case from Your Assessment

Your opportunity assessment is not just an exercise — it is the foundation of a business case. To convert your findings into a proposal that secures budget and executive support, structure it around four elements:

- **The problem:** 'Our team spends X hours per week on [task]. This costs $Y annually in labour and produces inconsistent quality.'
- **The solution:** 'AI tool [name] can automate [specific aspect] of this task, reducing time by [estimated percentage] while maintaining or improving quality.'
- **The evidence:** 'In our seven-day pilot, AI reduced task completion time from [X] minutes to [Y] minutes across [N] examples, with quality rated as [comparable/superior/requiring minor edits].'
- **The ask:** 'We recommend a ninety-day team-wide pilot at a cost of $[tool subscription] with expected savings of $[calculated from time reduction].'

10. Key Terms Glossary

Term	Definition
AI Opportunity Assessment	A structured evaluation of which tasks and processes in an organisation are best suited for AI automation or augmentation
Prioritisation Matrix	A two-dimensional framework (typically frequency vs. complexity) used to rank potential AI use cases by expected return and risk
Human-in-the-loop	A deployment model where AI performs a task but a human reviews, approves, or corrects the output before it is acted upon
Pilot project	A small-scale, time-limited test of a new technology on a real business task, designed to prove value before broader rollout
Error tolerance	The degree to which mistakes in a process can be absorbed without significant business, legal, or reputational consequences
ROI baseline	The measured current cost, time, and quality of a process before AI is introduced, used as the comparison point for calculating return on investment
Use case	A specific, defined application of technology to a particular business task or workflow
Stakeholder readiness	The willingness and ability of affected people to adopt and work with a new technology or process change

11. Quick Revision Checklist

- Start your assessment by listing the three most time-consuming repetitive tasks — specificity is critical
- Map each task to one of four AI capability categories: text generation, data analysis, research and synthesis, or creative content
- Prioritise using frequency (how often) and expertise level (how specialised) — high frequency, low-medium expertise tasks first
- Frequency is the ROI multiplier: a daily task saving thirty minutes delivers over one hundred hours of annual savings per person
- Lower-expertise tasks are safer first projects because errors are easier to catch
- Expand your assessment with data availability, process consistency, error tolerance, measurability, and stakeholder readiness
- Run a seven-day challenge: set up the tool, test three real examples, evaluate quality, document findings
- Personal experimentation converts more sceptics than any presentation or vendor demo
- Avoid starting too big (takes months) or too small (nobody cares about the result)
- Always measure the baseline before introducing AI — without it, you have anecdotes, not a business case
- Scale in concentric circles: same task to more people, then adjacent tasks, then different functions, then higher complexity

20 Exam-Based MCQs — AI Opportunity Assessment

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. ***Scenario:*** *A marketing director identifies three potential AI use cases: (1) generating weekly social media posts — done daily by a junior coordinator, (2) creating the annual brand strategy document — done once per year by the CMO, and (3) reformatting vendor invoices into the company's accounting template — done weekly by an accounts assistant.* Using the prioritisation matrix from this lecture, which task should be the FIRST AI pilot?

A) The annual brand strategy document, because it is the highest-value deliverable

B) The weekly social media posts, because they are the highest frequency and require moderate expertise

C) The vendor invoice reformatting, because it is the simplest task

D) All three should be piloted simultaneously to maximise learning

Answer: B **Explanation:** B is correct because daily frequency with moderate expertise requirements places social media posts squarely in Priority 1 of the matrix — high frequency, low-medium expertise. This delivers the fastest visible results with manageable risk. A is wrong because annual frequency and high expertise (CMO-level strategic work) places this in Priority 4 — complex and infrequent. C is wrong because while simplicity is appealing, weekly frequency is lower than daily, making the total time savings smaller; however, if the social media posts were not an option, this would be a reasonable Priority 1 alternative. D is wrong because simultaneous pilots split focus and make it harder to isolate what works; sequential pilots with clear measurement are more effective.

Q2. Why does the lecture emphasise specificity when listing time-consuming tasks?

A) Specific descriptions make the exercise more impressive to present to leadership

B) Vague descriptions prevent you from determining what AI needs to do, what inputs it requires, and how to measure success

C) AI tools require exact task descriptions to function properly

D) Specific tasks are always easier for AI to automate than general ones

Answer: B **Explanation:** B is correct because specificity transforms a vague aspiration ('automate reports') into an actionable scope ('compile monthly sales data from five sources into a PowerPoint template') — you can identify the inputs, the AI capability needed, and the success criteria. A is wrong because specificity serves operational clarity, not presentation quality. C is wrong because AI tools work with natural language prompts, not formal task descriptions; specificity helps your planning, not the tool's functionality. D is wrong because specificity helps you assess feasibility but does not inherently make a task easier to automate — a specific but highly complex task is still difficult.

Q3. ***Scenario:*** *A hospital administrator wants to use AI to improve patient scheduling. Currently, scheduling coordinators spend about forty percent of their time calling patients to confirm or reschedule appointments, twenty percent entering data into the EHR system, twenty percent handling insurance pre-authorisation checks, and twenty percent managing same-day cancellation backfills.* Which sub-task is the BEST candidate for an initial AI pilot?

A) Insurance pre-authorisation checks, because they involve the most regulatory complexity

B) Same-day cancellation backfills, because they are the most urgent

C) Patient confirmation and rescheduling calls, because they consume the most time and follow repetitive conversational patterns

D) EHR data entry, because it is the most structured task

Answer: C **Explanation:** C is correct because confirmation calls consume the largest share of time (forty percent), follow predictable conversational patterns, and can be handled by AI voice or text assistants with moderate complexity — fitting the high-frequency, low-medium expertise priority. A is wrong because insurance pre-authorisation involves regulatory complexity and specialist knowledge, making it a Priority 2 or 4 task depending on frequency. B is wrong because urgency does not determine pilot suitability; backfills are time-sensitive but may require human judgment about patient priority. D is wrong because while structured data entry could be automated, it consumes only twenty percent of time — half the impact of the confirmation calls.

Q4. A task that is performed monthly and requires deep specialist knowledge falls into which priority category?

A) Priority 1 — start here for quick wins

B) Priority 2 — high value but needs expert oversight

C) Priority 3 — moderate impact, low setup effort

D) Priority 4 — park this until the team has AI experience

Answer: D **Explanation:** D is correct because the prioritisation matrix places low-frequency (monthly), high-expertise tasks in Priority 4. These tasks are poor first AI projects because they happen too infrequently to generate visible momentum and require specialist oversight that the organisation may not yet have established for AI outputs. A is wrong because Priority 1 is for high-frequency, low-medium expertise tasks. B is wrong because Priority 2 is for high-frequency, high-expertise tasks — the frequency makes them worth the oversight investment. C is wrong because Priority 3 is for low-frequency, low-medium expertise tasks.

Q5. ***Scenario:*** *A VP of Sales introduces ChatGPT to the sales team for drafting prospecting emails. After two weeks, the team reports that the emails are good but not great — they need about five minutes of editing per email to match the company's tone and add specific product details. Previously, writing each email from scratch took about twenty minutes.* How should the VP interpret this result?

A) The AI is not good enough and should be abandoned for a better tool

B) The AI reduced a twenty-minute task to five minutes of editing — a seventy-five percent time saving that represents a successful pilot worth scaling

C) The team should stop using AI until it can produce perfect emails without editing

D) The AI is underperforming because it should eliminate the need for human involvement entirely

Answer: B **Explanation:** B is correct because a seventy-five percent reduction in time per task (from twenty minutes to five minutes of editing) is a substantial productivity gain. AI does not need to be perfect to be valuable — it needs to be better than the current process. A is wrong because the AI is delivering clear value; switching tools based on a need for minor editing is premature optimisation. C is wrong because waiting for perfection means waiting indefinitely; the current improvement is already delivering ROI. D is wrong because human-in-the-loop editing is the expected model for AI-assisted content creation, not a failure state.

Q6. Which evaluation criterion helps determine whether an AI use case is actually a data engineering project in disguise?

A) Error tolerance

B) Stakeholder readiness

C) Data availability

D) Process consistency

Answer: C **Explanation:** C is correct because data availability reveals whether the inputs AI needs are accessible and digital. If the required data is scattered across disconnected systems, stored in paper files, or locked in formats AI cannot process, the project is primarily a data integration effort that must be completed before AI can add value. A is wrong because error tolerance measures risk appetite, not data readiness. B is wrong because stakeholder readiness addresses people and change management, not technical prerequisites. D is wrong because process consistency evaluates how predictable the task is, not whether its inputs are available.

Q7. ***Scenario:*** *An operations manager completes the AI opportunity assessment and identifies email triage as the top priority — her team of eight people each spends ninety minutes per day reading, categorising, and routing incoming supplier emails. She wants to build a business case for her COO.* What is the MOST compelling way to frame the business case?

A) 'AI is the future and we need to adopt it to stay competitive'

B) 'Our team spends 960 hours per month on email triage. A ninety-day pilot with an AI tool at $200 per month could reduce this by sixty percent, saving 576 hours monthly — equivalent to 3.6 full-time employees'

C) 'ChatGPT can read emails, so we should buy licenses for everyone'

D) 'Our competitors are using AI for email management'

Answer: B **Explanation:** B is correct because it follows the business case structure: quantified problem (960 hours/month), specific solution (AI email triage), evidence-based projection (sixty percent reduction), and financial translation (3.6 FTE equivalent) — exactly the kind of data-driven argument a COO needs to approve budget. A is wrong because 'the future' is not a business case; it is a platitude that does not help a COO allocate resources. C is wrong because it jumps to a solution without defining the problem, quantifying the opportunity, or proposing measurement criteria. D is wrong because competitive pressure alone does not quantify the internal benefit or justify the specific investment.

Q8. What is the PRIMARY purpose of measuring a baseline before introducing AI to a task?

A) To prove that the current process is inefficient enough to justify AI investment

B) To provide a comparison point so you can objectively calculate the improvement AI delivers

C) To document the process for the AI training data

D) To identify which employees are underperforming

Answer: B **Explanation:** B is correct because without a pre-AI measurement of time, cost, and quality, you cannot calculate ROI. The baseline converts 'AI seems helpful' into 'AI reduced task time by X percent and maintained Y quality score.' A is wrong because the baseline's purpose is measurement, not justification; you should already believe the task is worth improving before measuring it. C is wrong because baseline measurement captures process metrics, not training data for the AI model. D is wrong because baseline measurement is about the process, not individual performance evaluation.

Q9. ***Scenario:*** *A financial services firm's compliance officer identifies AI-assisted regulatory document review as a high-priority opportunity. The task is performed daily, takes four hours per analyst, and the team has six analysts. However, regulatory review errors could result in fines of up to five million dollars.* How should the compliance officer position this in the prioritisation matrix?

A) Priority 1 — high frequency and clear time savings make it a quick win

B) Priority 2 — high frequency justifies the investment, but high expertise and error consequences require a human-in-the-loop design with rigorous testing before rollout

C) Priority 4 — the risk is too high for any AI involvement

D) Priority 3 — the risk makes it a lower priority regardless of frequency

Answer: B **Explanation:** B is correct because the task is high-frequency (daily) and high-expertise (regulatory knowledge with significant error consequences), placing it in Priority 2. The frequency makes the potential savings enormous (24 analyst-hours per day), justifying the investment in human-in-the-loop governance. A is wrong because it ignores the high expertise and severe error consequences; this is not a 'quick win' but requires careful governance design. C is wrong because high frequency and massive time savings make this worth pursuing — with appropriate safeguards, not avoidance. D is wrong because Priority 3 is for low-frequency, low-expertise tasks; this task is high-frequency and high-expertise.

Q10. When scaling AI beyond the first successful pilot, what is the recommended FIRST expansion step?

A) Apply AI to the most complex task in the organisation to demonstrate maximum impact
B) Roll out the same AI application to more people performing the same task
C) Implement AI in a completely different department to prove cross-functional value
D) Invest in building a custom AI model tailored to the organisation

Answer: B **Explanation:** B is correct because same task, more people is the lowest-risk expansion — the use case is proven, the measurement framework exists, and the training approach is established. You are scaling something that works, not experimenting with something new. A is wrong because jumping to maximum complexity before building organisational AI capability is the 'starting too big' pitfall. C is wrong because while cross-functional expansion is valuable, it is the third step, not the first — you should first deepen usage within the proven function. D is wrong because custom model development is expensive and premature before the organisation has validated that AI delivers value on standard tools.

Q11. ***Scenario:*** *A chief human resources officer has identified three potential AI use cases: (1) screening the first round of job applications by matching qualifications against role requirements — done fifty times per week; (2) writing personalised onboarding welcome emails — done five times per week; (3) analysing exit interview transcripts for common themes — done quarterly.* What is the correct priority ordering?

A) 1 (screening), then 2 (onboarding emails), then 3 (exit interviews)
B) 3 (exit interviews), then 1 (screening), then 2 (onboarding emails)
C) 2 (onboarding emails), then 1 (screening), then 3 (exit interviews)
D) All three should be implemented simultaneously

Answer: A **Explanation:** A is correct because application screening is the highest frequency (fifty per week) with moderate expertise; onboarding emails are lower frequency (five per week) but straightforward; exit interview analysis is quarterly (lowest frequency). The matrix prioritises by frequency first, with expertise as the secondary factor. B is wrong because quarterly tasks should not be prioritised over weekly ones. C is wrong because onboarding emails at five per week are lower frequency than screening at fifty per week. D is wrong because sequential implementation allows learning and resource focus.

Q12. A manager says: 'We tried AI on our expense reports and it saved each person about two minutes per report. That proves AI works — let's roll it out company-wide.' What is the MAIN risk with this reasoning?

A) Two minutes per report is too small a saving to justify company-wide rollout costs and change management effort

B) AI should not be used for financial tasks under any circumstances

C) The manager should have tested with a larger sample size first

D) Expense reports are too complex for AI

Answer: A **Explanation:** A is correct because two minutes per report — even at scale — may not justify the cost of tool licences, training, and change management for a company-wide rollout. This is the 'choosing a task nobody cares about' pitfall: the saving is real but insufficient to generate organisational enthusiasm or meaningful ROI. B is wrong because AI is used successfully in many financial tasks; the issue is impact size, not task category. C is wrong because sample size is a valid consideration but secondary to the fundamental question of whether the per-instance savings justify the rollout investment. D is wrong because the AI clearly worked; the issue is whether the impact is large enough to scale, not whether the task is feasible.

Q13. Which of the following is the BEST example of a well-specified task for an AI opportunity assessment?

A) 'Improve customer service'

B) 'Drafting the weekly three-paragraph project status update email to stakeholders from Jira ticket data and team stand-up notes'

C) 'Use AI somewhere in marketing'

D) 'Make the team more productive'

Answer: B **Explanation:** B is correct because it specifies exactly what is being produced (a three-paragraph status email), how often (weekly), the inputs required (Jira data and stand-up notes), and the recipients (stakeholders). This level of detail allows you to identify the AI capability (text generation + data synthesis), estimate time savings, and design a pilot. A, C, and D are all too vague to act on — they describe aspirations, not automatable tasks.

Q14. ***Scenario:*** *An IT director completes an AI opportunity assessment and identifies help desk ticket categorisation as the top priority. The task is done two hundred times per day and currently takes an average of three minutes per ticket. Before proceeding, a colleague warns: 'Our ticket data is stored in three different systems that don't talk to each other, and thirty percent of tickets arrive by phone call with no written record.'* What should the IT director conclude?

A) AI cannot be used for help desk operations and the project should be abandoned

B) The AI project is actually a data integration project first — the three systems need consolidation and phone tickets need transcription before AI can effectively categorise at scale

C) The IT director should proceed with AI on the seventy percent of tickets that are already digital

D) A different AI tool that can handle disconnected data sources should be selected

Answer: B **Explanation:** B is correct because the data availability assessment revealed that the prerequisite infrastructure (unified ticketing data, transcribed phone calls) does not exist. This is the 'confusing tool capability with business readiness' pitfall — AI can categorise tickets, but only if it can access them. The project scope must start with data integration. C is tempting but wrong as a primary conclusion because ignoring thirty percent of volume creates an incomplete system and does not address the underlying fragmentation. A is wrong because the opportunity is real; it just requires prerequisite work. D is wrong because no AI tool can magically integrate disconnected backend systems.

Q15. According to the lecture, what is the MOST effective way to convert AI sceptics in an organisation?

A) Sharing industry research and analyst reports about AI's potential

B) Having sceptics try AI on their own real tasks during a structured personal experimentation period

C) Presenting case studies from other companies

D) Mandating AI usage through a company-wide policy

Answer: B **Explanation:** B is correct because the lecture emphasises that personal experimentation — trying AI on your own tasks and experiencing the results firsthand — converts sceptics more effectively than any presentation, report, or case study. People trust their own experience over someone else's data. A is wrong because research reports are informative but abstract; they do not create the personal 'aha' moment that drives adoption. C is wrong because case studies from other companies feel distant — 'that worked for them, but our situation is different.' D is wrong because mandating usage creates compliance, not enthusiasm, and often triggers resistance.

Q16. ***Scenario:*** *After a successful seven-day AI pilot on sales email drafting, a VP of Sales wants to expand. The sales team has twelve people. The VP is considering three next steps: (1) Roll out email drafting to all twelve sales reps. (2) Try AI for competitive intelligence research for the sales team. (3) Introduce AI email drafting to the customer success team.* What is the recommended sequence?

A) 2, then 3, then 1

B) 1, then 2, then 3

C) 3, then 1, then 2

D) 1, then 3, then 2

Answer: B **Explanation:** B is correct because it follows the scaling sequence: (1) same task, more people — proven use case, lowest risk; (2) adjacent task, same function — expands AI within sales using the team's growing AI fluency; (3) same capability, different function — takes the proven email drafting capability to customer success. A is wrong because it skips the lowest-risk step (scaling the proven use case). C is wrong because expanding to a different function before scaling within the proven team is premature. D is wrong because going to customer success (different function) before trying adjacent tasks within sales skips a step.

Q17. What distinguishes a Priority 2 task from a Priority 1 task in the opportunity assessment matrix?

A) Priority 2 tasks are less frequent than Priority 1 tasks

B) Priority 2 tasks require higher expertise and therefore need expert-in-the-loop oversight, even though they are high-frequency

C) Priority 2 tasks have lower business impact

D) Priority 2 tasks cannot be automated with current AI technology

Answer: B **Explanation:** B is correct because both Priority 1 and Priority 2 tasks are high-frequency; the difference is expertise level. Priority 1 tasks require low-medium expertise (minimal oversight), while Priority 2 tasks require high expertise (specialist review of AI outputs). The higher oversight cost is justified by the high frequency. A is wrong because both priorities share the high-frequency characteristic. C is wrong because Priority 2 tasks can have very high business impact — their ranking is about implementation sequencing, not value. D is wrong because Priority 2 tasks are automatable; they just need more careful governance.

Q18. ***Scenario:*** *A small business owner asks her accountant to evaluate AI for the business. The accountant identifies that the owner spends three hours every Monday compiling a cash flow summary from bank statements, invoices, and receipt scans. The bank statements are digital, the invoices are in QuickBooks, but the receipts are paper slips kept in a shoebox.* What is the CORRECT assessment?

A) AI can automate the entire cash flow compilation immediately

B) AI can process the digital bank statements and QuickBooks invoices now, but the paper receipts need to be digitised (scanned or photographed) before AI can include them — a partial automation is possible immediately

C) The task cannot use AI until all inputs are in the same system

D) Paper receipts are too complex for AI to ever process

Answer: B **Explanation:** B is correct because it applies the data availability criterion practically: two of three data sources are already digital and can be processed by AI now, while the third (paper receipts) requires a digitisation step. A pragmatic approach is to automate what is immediately possible and address the gap. A is wrong because it ignores the data availability constraint of the paper receipts. C is wrong because partial automation is valuable — processing two of three sources automatically still saves significant time. D is wrong because AI can process scanned or photographed receipts (OCR + interpretation); the receipts just need to be digitised first.

Q19. Why does the lecture recommend starting with three real examples during the seven-day challenge rather than a single test?

A) Three examples are the minimum required for statistical significance

B) Three examples provide enough variation to see how AI handles different instances of the same task, giving a more realistic assessment of its capabilities and limitations

C) AI tools require at least three examples to calibrate their outputs

D) Three is the maximum number of free uses most AI tools allow

Answer: B **Explanation:** B is correct because a single example might be unusually easy or unusually hard for AI, giving a skewed impression. Three examples with natural variation (different subjects, different complexity levels) provide a more representative picture of what to expect at scale. A is wrong because this is a practical pilot, not a scientific study; three examples is about practical assessment, not statistical rigour. C is wrong because AI tools do not need calibration examples from users. D is wrong because most AI tools offer more than three free uses and pricing is not the reason for the recommendation.

Q20. ***Scenario:*** *A department head runs a successful seven-day AI pilot on report drafting. She saved an average of forty-five minutes per report across her three test reports. She now needs to present findings to the executive committee. She has data on time saved but did not assess the quality of the AI-generated reports compared to manually written ones.* What is the MOST significant gap in her pilot evaluation?

A) She should have tested with more than three reports

B) She did not compare the quality of AI outputs against manually written reports, making her business case incomplete — time savings without quality validation could mean the team spends the 'saved' time on rework

C) She should have used multiple AI tools to compare performance

D) She did not get executive approval before running the pilot

Answer: B **Explanation:** B is correct because time savings are meaningless if quality drops so far that the 'saved' time is consumed by editing and rework. A complete pilot evaluation measures both efficiency (time) and effectiveness (quality) to present a credible business case. A is wrong because while more examples improve confidence, three is sufficient for an initial assessment — the missing quality comparison is a more fundamental gap. C is wrong because comparing tools is a secondary optimisation step; proving the concept works at all comes first. D is wrong because the purpose of a seven-day pilot is lightweight experimentation that typically does not require formal executive approval.